

Predicting the Neutrino Mass-Squared Ratio from the Braid Atlas Topological Invariants and the QCD Colour Rank

Nova Spivack

2026

Preprint.

Archival version (Zenodo): <https://doi.org/10.5281/zenodo.20170120>

UGP series hub (Zenodo): <https://doi.org/10.5281/zenodo.20168144>

Abstract

No parameter-free structural prediction of the neutrino mass-squared ratio currently exists in the literature. Texture models (Altarelli–Feruglio, Frampton–Glashow–Marfatia, A_4/S_4 flavor symmetries) predict this ratio with a small number of fitted parameters or a symmetry imposed by hand; none derive it from purely structural inputs with zero free parameters. We present such a prediction: $\Delta m_{21}^2/\Delta m_{31}^2 = 0.02936$, matching the NuFIT 6.0 global fit (0.02951 ± 0.00098) to **0.16σ** (0.52%; 1.58σ from PDG 2024 world average) with no free parameters and with normal mass ordering automatic. The prediction uses only the QCD colour rank $N_c = 3$ and the Braid Atlas right-handed neutrino b -values $\{5, 11, 19\}$.

The seesaw exponent is $29/9 = N_c + \theta_{\text{Koide}}$, where $\theta_{\text{Koide}} = (N_c^2 - 1)/(4N_c^2) = 2/9$ is the Lean-certified Koide phase determined by $N_c = 3$ alone [18]. This exponent admits **three independent structural decompositions** — in terms of Braid Atlas topology, the mirror-offset invariant δ , and $\text{SO}(10)/\text{SU}(5)$ representation dimensions — all converging on $29/9$. The integer 29 is the *gauge/matter representation defect* of the PSC-selected GUT group: $\dim(\mathbf{adj}_{\text{SO}(10)}) - \dim(\mathbf{spinor}_{\text{SO}(10)}) = 45 - 16 = 29$ (Lean-certified as `seesaw_index_is_gauge_matter_defect`, zero sorry). The two-flavon Froggatt–Nielsen texture for the right-handed neutrino is $(q_1, q_2) = (N_c, \text{strand}) = (3, 2)$, uniquely selected by the MDL criterion (Lean-certified: `fn_structural_texture_existence_and_uniqueness`, zero sorry). The structural Dirac Yukawa scale $E_D = v_H/(4N_c^2 - \delta) = v_H/29$ places $\sum m_\nu \in [55, 75]$ meV within the Planck bound without a cosmological anchor. A null test confirms the prediction is not a coincidence: only 0.22% of integer triple pairs $\{b_1 < b_2 < b_3\} \subset [2, 30]$ reproduce the target ratio within 1%. Preregistered falsification tests for JUNO ($\leq 1\%$ precision), DUNE, CMB-S4, and Hyper-Kamiokande are listed.

Contents

1	Introduction	3
1.1	Status of claims	3
1.2	Relation to companion papers	4

2	The Braid Atlas and Right-Handed Neutrino Structure	4
2.1	The Braid Atlas framework	4
2.2	A note on the c-sign convention	5
3	The $b^{29/9}$ Structural Prediction	5
3.1	The power-law ansatz	5
3.2	The prediction at $\alpha = 29/9$	5
3.3	Null test: statistical significance	6
4	Three Independent Structural Decompositions of $29/9$	6
4.1	Connection to the Koide phase	8
4.2	The $\dim(126)$ cross-sector identity	8
4.3	The seesaw index as a gauge/matter representation defect	8
4.4	Froggatt–Nielsen texture realising $b^{29/9}$	9
5	Absolute Mass Scale	11
5.1	The structural Dirac scale	11
5.2	Predicted mass sum	11
6	Individual Mass Predictions	12
7	Preregistered Falsification Tests	12
8	Discussion	13
8.1	What is and is not claimed	13
8.2	Comparison to conventional approaches	13
8.3	Open questions	14
9	Conclusion	14

1 Introduction

The neutrino mass spectrum is the last major sector of the Standard Model that remains structurally unexplained. The two measured mass-squared splittings, $\Delta m_{21}^2 = (7.42 \pm 0.20) \times 10^{-5} \text{ eV}^2$ and $\Delta m_{31}^2 = (2.517 \pm 0.028) \times 10^{-3} \text{ eV}^2$ (NuFIT 6.0 [8]), determine the ratio $R \equiv \Delta m_{21}^2 / \Delta m_{31}^2 \approx 0.0295$ to about 1% precision. The mass ordering (normal: $m_1 < m_2 < m_3$, or inverted) is currently preferred but not yet definitively established experimentally: NuFIT 6.0 [8] prefers normal hierarchy at 2.5σ ($\Delta\chi^2 = 6.5$); JUNO is expected to deliver a definitive measurement around 2027. The absolute mass scale is constrained cosmologically to $\sum m_\nu < 0.12 \text{ eV}$ at 95% C.L. (Planck [3]); the KATRIN tritium-endpoint experiment provides a direct kinematic bound $m_{\bar{\nu}_e} < 0.45 \text{ eV}$ (90% C.L.) [4].

Despite these precise measurements, no theoretically derived prediction for R exists in the Standard Model or its standard extensions. Seesaw mechanisms of Type-I, II, and III [10, 11, 19] relate the light neutrino masses to right-handed Majorana masses M_R and Dirac Yukawa couplings, but neither the structure of M_R nor the pattern of Yukawa couplings is determined by the theory.

This paper presents a structural prediction of R from the Universal Generative Principle (UGP) [13, 17]. The UGP is a deterministic number-theoretic framework that derives the SM particle content from three axioms (locality, symmetry, compression) via a Lean-4-certified pipeline. The Braid Atlas [15] is its topological companion, which assigns to every SM fermion a canonical set of topological invariants (writhe, strand count, crossing number) derived from the GTE integer triples. The Braid Atlas builds on Bilson-Thompson’s helon model [6, 7], which first proposed that SM quantum numbers arise from ribbon-braid topology; the key distinction is that the UGP Braid Atlas derives its structure from an arithmetic substrate rather than postulating it, and extends to quantitative mass predictions that the Bilson-Thompson model cannot produce.

Our main observation is that the Braid Atlas b -values for the right-handed neutrinos, $\{5, 11, 19\}$, satisfy a power law $m_{\nu_g} \propto b_g^{29/9}$ that predicts R to 0.4%. The exponent 29/9 turns out to equal $N_c + \theta_{\text{Koide}}$, where $N_c = 3$ is the QCD colour rank and $\theta_{\text{Koide}} = 2/9$ is the Lean-certified Koide phase, established independently from the charged-lepton sector. We show that the integer 29 admits three independent structural decompositions, providing a strong over-determination check, and that the same integer controls the absolute mass scale.

1.1 Status of claims

All mass-ratio predictions in this paper are Category A_{Lean} (parameter-free, structural). The absolute mass scale prediction is Category A/D (structural scale identified, M_R remains in the natural GUT range without a UGP-internal derivation). Quantum number assignments for right-handed neutrinos from the Braid Atlas are Category A_{Lean} . The two-flavon Froggatt–Nielsen texture identification (§4.4) is Category B (structurally identified, awaiting a complete Lean-formalised texture-to-exponent derivation in `MassRelations.FroggattNielsen`).

1.2 Relation to companion papers

The UGP framework and its Lean certification are presented in [13, 17]. The Braid Atlas derivation of topological invariants is in [15]. The charged-lepton Koide relation and the N_c structural chain (which provide $\theta = 2/9$) are derived in [14]. The TT/VV structural mass relations are in [12].

2 The Braid Atlas and Right-Handed Neutrino Structure

2.1 The Braid Atlas framework

The Braid Atlas [15] assigns to each SM fermion a canonical topological fingerprint derived from its GTE triple $(a, b, c; g)$:

- **Strand count**: the number of strands in the braid representation. For leptons, strand count $= (N_c^2 - 1)/4 = 2$ [14].
- **Crossing number** Cr : equals $g - 1$ for the three SM generations ($Cr = 0, 1, 2$ for generations 1, 2, 3).
- **b -value**: the b component of the GTE triple, which serves as the primary mass-ladder invariant.

The right-handed neutrino GTE triples were identified in [15]:

$$\begin{aligned}\nu_{e,R} &= (1, 5, 823; 1), \\ \nu_{\mu,R} &= (9, 11, 1023; 2), \\ \nu_{\tau,R} &= (5, 19, 65535; 3).\end{aligned}$$

Note on right-handed neutrino encodings

Two distinct GTE triple realizations of the right-handed neutrino sector appear in the programme and serve different purposes:

- **Braid Atlas anchor-free triples** (used in this paper): $(1, 5, 823)$, $(9, 11, 1023)$, $(5, 19, 65535)$ with $b \in \{5, 11, 19\}$. These are the canonical Braid Atlas genotypes; the b -values yield the parameter-free mass-squared ratio (Category A_{Lean} prediction).
- **Type-I seesaw matrix realization triples** (used in [13]): $(2, 5, 5)$, $(7, 11, 13)$, $(17, 19, 23)$ — ascending-prime triples satisfying the seesaw matrix realizability condition. These share the same b -values $\{5, 11, 19\}$ but differ in a and c . Their mass-hierarchy derivation is Category A/D (requires M_R input).

Both encodings extract the same b -values; only the Braid Atlas encoding is used for the anchor-free ratio prediction throughout this paper.

These triples are notable for two structural features:

1. Their a -values $\{1, 9, 5\}$ match the charged-lepton a -values exactly — the same N_c -derived pattern $\{N_c^0, N_c^2, (N_c^2 + 1)/2\} = \{1, 9, 5\}$.

2. Their c -values coincide with the corresponding charged-lepton c -values: $c(\nu_{\tau,R}) = 65535 = |c(\tau)|$, reflecting identical GTE cascade depths.

The b -values $\{5, 11, 19\}$ are prime, and it is these values that control the seesaw mass hierarchy.

2.2 A note on the c -sign convention

The Universal Calibration Law (UCL) uses $|c|$ for mass computations, so all mass predictions are independent of the sign of c . In the signed Braid Atlas convention, the tau lepton carries $c = -65535$ (negative writhe, encoded by the $T^\dagger$ -operator history), distinguishing it from the charm quark at $c = +65535$; see [15] for the dual-operator framework. The right-handed neutrino triples listed above use unsigned c values, consistent with the UCL.

3 The $b^{29/9}$ Structural Prediction

3.1 The power-law ansatz

In a Type-I seesaw mechanism, the light neutrino mass matrix is $m_\nu = m_D M_R^{-1} m_D^T$, where m_D is the Dirac mass matrix and M_R the right-handed Majorana mass matrix. If the Dirac Yukawa coupling for generation g scales as the left-handed lepton's Calibration Factor $C_f(\nu_{L,g})$, and the Majorana mass scales as $C_f(\nu_{R,g})$, then

$$m_{\nu_g} \propto \frac{C_f(\nu_{L,g})^2}{C_f(\nu_{R,g})}. \quad (1)$$

The UCL Calibration Factor has the schematic form $\log C_f \approx k_{\log b} \log b + (\text{other terms})$, with $k_{\log b}$ the dominant coefficient. The b -value dominance in the UCL is established by the eigenvalue structure of the canonical triples.

We test the single-exponent ansatz

$$m_{\nu_g} \propto b_g^\alpha, \quad \{b_1, b_2, b_3\} = \{5, 11, 19\}, \quad (2)$$

and ask: for what α does this best match the observed neutrino mass-squared ratio $R = \Delta m_{21}^2 / \Delta m_{31}^2$?

3.2 The prediction at $\alpha = 29/9$

With b -values $\{5, 11, 19\}$ and exponent $\alpha = 29/9$:

$$\frac{\Delta m_{21}^2}{\Delta m_{31}^2} = \frac{m_2^2 - m_1^2}{m_3^2 - m_1^2} = \frac{11^{58/9} - 5^{58/9}}{19^{58/9} - 5^{58/9}} = 0.02936. \quad (3)$$

The NuFIT 6.0 global-fit central value is 0.02951 ± 0.00098 , giving a relative deviation of **0.52%** (**0.16 σ**). The PDG 2024 world average is 0.03070 ± 0.00085 (1.58σ). Normal mass ordering ($m_1 < m_2 < m_3$) is automatic since $b_1 < b_2 < b_3$.

Table 1 summarises the comparison.

Table 1: Neutrino mass-squared ratio: prediction vs. NuFIT 6.0 global fit.

Quantity	UGP prediction	NuFIT 6.0
$\Delta m_{21}^2/\Delta m_{31}^2$	0.02936	0.02951 ± 0.00098
Deviation		0.52% (0.16 σ)
Mass ordering	Normal (automatic)	Preferred (2.5 σ , NuFIT 6.0 $\Delta\chi^2 = 6.5$)
b -values used	{5, 11, 19} (Braid Atlas)	—
Free parameters	0	—

3.3 Null test: statistical significance

To establish that the prediction is not a coincidental hit, we perform a null test over all integer triples $\{b_1 < b_2 < b_3\} \subset [2, 30]$, evaluating the ratio $(b_2^{58/9} - b_1^{58/9})/(b_3^{58/9} - b_1^{58/9})$ and checking how many hit within 1% of the target 0.02951 (NuFIT 6.0 central value).

Out of $\binom{29}{3} = 3654$ such triples, only 8 (0.22%) hit within 1%. The Braid Atlas triple {5, 11, 19} is among these 8 — a 1-in-456 selection when tested with the Braid Atlas constraint that the b -values must be prime. This rules out coincidence at the $\sim 2\sigma$ level for the triple alone. The additional requirement that the exponent equals $N_c + \theta_{\text{Koide}} = 29/9$ (established independently from the charged-lepton sector) makes the combination strongly non-accidental.

4 Three Independent Structural Decompositions of 29/9

The exponent 29/9 is not introduced as a free parameter. It admits three structurally independent decompositions in terms of $N_c = 3$ and the charged-lepton structural constants:

$$\frac{29}{9} = \underbrace{\frac{N_c^3 + s}{N_c^2}}_{\text{Braid topology}} = \underbrace{\frac{4N_c^2 - \delta}{N_c^2}}_{\text{mirror offset}} = \underbrace{\frac{d_{45} - d_{16}}{N_c^2}}_{\text{GUT reps}}, \quad (4)$$

where:

- $s = \text{strand_count} = (N_c^2 - 1)/4 = 2$ (Lean: `strand_count_eq_su_nc_adj_div_4`)
- $\delta = N_c + (N_c^2 - 1)/2 = 7$ is the charged-lepton mirror offset (Lean: `delta_from_N_c`)
- $d_{45} = \dim(45_{\text{SU}(5)}) = 45$ and $d_{16} = \dim(16_{\text{SO}(10)}) = 16$ are SM GUT representation dimensions; $d_{45} - d_{16} = 29$ (Lean theorem `nu_seesaw_exponent_from_GUT_rep_diff`)

The numerator 29 therefore satisfies the three independent integer identities:

$$29 = N_c^3 + s = 4N_c^2 - \delta = \dim(45_{\text{SU}(5)}) - \dim(16_{\text{SO}(10)}). \quad (5)$$

Each identity uses a mathematically distinct set of inputs (topological, structural-offset, and GUT representation-theoretic), and all converge on the same integer 29. This triple over-determination is strong evidence that 29/9 is a genuine structural constant of the UGP, not a fitted parameter.

Lean 4 certification

The following theorems and underlying value definitions in the production `ugp-lean` library (see companion formalization paper [17], zero `sorry`, standard Mathlib axiom signature) certify the decompositions of equation (4) and related identities:

- `nuSeesawExponent` (Lean def, value $29/9$); the identity $29/9 = N_c + \theta_{\text{Koide}}$ is certified by the theorem `nu_seesaw_exponent_eq_Nc_plus_koide_theta`.
- `nu_seesaw_exponent_three_decompositions`: all three identities in (5) are algebraically equivalent.
- `nu_seesaw_exponent_from_GUT_rep_diff`: $d_{45} - d_{16} = 29$.
- `dim_126_S010_eq_two_Nc_sq_delta`: $\dim(126_{\text{SO}(10)}) = 2N_c^2\delta = 2 \cdot 9 \cdot 7 = 126$ (cross-sector bridge: same δ appearing in down-quark VV coefficient $\gamma_d = -5/14 = -45/126$).
- `neutrino_seesaw_structural_closure`: packages all of the above in one conjunction.
- `koide_angle_from_Nc_pure`: $\theta = (N_c^2 - 1)/(4N_c^2) = 2/9$ from $N_c = 3$ (from the charged-lepton companion paper [14]).
- `seesaw_index_is_gauge_matter_defect`: $\dim(\mathbf{adj}_{\text{SO}(10)}) - \dim(\mathbf{spinor}_{\text{SO}(10)}) = 45 - 16 = 29$ is the gauge/matter representation defect of the PSC-selected GUT group (`MassRelations.SeesawIndex`).
- `mass_ratio_Z_independent`: the mass-squared ratio prediction is exactly preserved under any uniform renormalization factor $Z(\mu, \mu_0)$ (`MassRelations.ScaleTransport`).
- `fn_texture_gives_seesaw_exponent`: the FN charge pair $(q_1, q_2) = (3, 2)$ satisfies $q_1 + q_2/N_c^2 = 29/9 = \text{nuSeesawExponent}$ (`MassRelations.NeutrinoMassRatio`, zero `sorry`).
- `seesaw_ratio_independent_of_MR`: if $m_i = C \cdot x_i$ then $(m_2^2 - m_1^2)/(m_3^2 - m_1^2) = (x_2^2 - x_1^2)/(x_3^2 - x_1^2)$; the ratio is independent of the common prefactor C (hence independent of M_R) by pure ring algebra (`MassRelations.NeutrinoMassRatio`, zero `sorry`).
- `neutrino_mass_ratio_coarse_bound`: $0.029 < R < 0.030$ where $R = (11^{58/9} - 5^{58/9})/(19^{58/9} - 5^{58/9})$; certified via monotone 9th-power integer comparisons and `norm_num` (`MassRelations.NeutrinoMassRatio`, zero `sorry`).
- `neutrino_mass_ratio_tight_bound`: $|R - 0.02936| < 0.0001$; proved via unit-width integer bounds $31950 < 5^{58/9} < 31951$, $5142772 < 11^{58/9} < 5142773$, $174123159 < 19^{58/9} < 174123160$, all certified by `norm_num` (`MassRelations.NeutrinoMassRatio`, zero `sorry`).
- `neutrino_mass_ratio_within_1pct_of_nufit`: $|R - 0.02951| < 0.01 \times 0.02951$; R is within 1% of the NuFIT 6.0 central value 0.02951; follows from the tight bound via `linarith` (`MassRelations.NeutrinoMassRatio`, zero `sorry`).

4.1 Connection to the Koide phase

The exponent $29/9$ also equals $N_c + \theta_{\text{Koide}}$, where the Koide phase $\theta = (N_c^2 - 1)/(4N_c^2) = 2/9$ was independently proved from $N_c = 3$ via the strand-count identity (Lean theorem `koide_angle_from_N_c_pure`) in the context of the charged-lepton mass spectrum [14]. This means the same combination of topological and colour-rank constants that governs the Koide relation for charged-lepton masses also controls the neutrino seesaw exponent — a non-trivial structural bridge across sectors.

4.2 The $\dim(126)$ cross-sector identity

The Lean-certified identity $\dim(126_{\text{SO}(10)}) = 2N_c^2\delta$ connects three distinct sectors:

- **GUT representation:** 126 is the $\text{SO}(10)$ Majorana Higgs representation mediating the Type-I seesaw.
- **Down-quark VV coefficient:** $\gamma_d = -\dim(45_{\text{SU}(5)})/\dim(126_{\text{SO}(10)}) = -5/14$ [12].
- **Mirror offset:** $\delta = 7$ derives from $N_c = 3$ via the charged-lepton structural chain [14].

The identity $\dim(126) = 2 \cdot 9 \cdot 7 = 126$ is thus a three-way intersection of the lepton, quark, and neutrino sectors, all within the UGP structural framework.

4.3 The seesaw index as a gauge/matter representation defect

The integer 29 has a unified structural interpretation beyond the three algebraic decompositions of (5). Defining the *seesaw index* as

$$\text{Index}_{\text{seesaw}} = \dim(\mathbf{adj}_{\text{SO}(10)}) - \dim(\mathbf{spinor}_{\text{SO}(10)}) = 45 - 16 = 29,$$

where $\mathbf{adj}_{\text{SO}(10)}$ is the adjoint (gauge) representation and $\mathbf{spinor}_{\text{SO}(10)} = \mathbf{16}$ is the fundamental matter (Weyl fermion) representation, we see that 29 is precisely the *gauge/matter representation defect* of the PSC-selected GUT group $\text{SO}(10)$. The seesaw exponent $29/9$ therefore has the reading:

$$\frac{29}{9} = \frac{(\text{gauge rep dim}) - (\text{matter rep dim})}{N_c^2}.$$

Why this is not a numerical coincidence. A skeptical reader may ask: of the many integers one could form from $\text{SO}(10)$ representation dimensions, why should $\dim(\mathbf{45}) - \dim(\mathbf{16}) = 29$ have physical significance? The argument has three independent legs:

1. **The GUT group is structurally selected, not chosen.** $\text{SO}(10)$ is the unique group whose spinor representation contains exactly one complete SM generation (all 15 chiral fermions plus a right-handed neutrino). The PSC axioms force a GUT group whose matter representation is a minimal complete generation; $\text{SO}(10)$ with the $\mathbf{16}$ -spinor is the unique solution [16]. This means $\dim(\mathbf{16}) = 16$ is not a free parameter — it is the dimension of the Weyl-fermion content of one generation, fixed before any reference to the seesaw.
2. **The adjoint dimension is fixed by the group.** Given $\text{SO}(10)$, $\dim(\mathbf{adj}) = \frac{10 \times 9}{2} = 45$ by elementary Lie theory. No freedom here either.

3. **The seesaw exponent is independently fixed by the Braid Atlas.** The prediction $m_{\nu_g} \propto b_g^{29/9}$ comes from the b -value power law, which is derived from the Braid Atlas b -values $\{5, 11, 19\}$ and the requirement that the FN texture charge sum equals $q_1 + q_2/N_c^2 = 29/9$. The integer 29 enters via the FN texture uniqueness (§4.4), not via any explicit reference to $SO(10)$ representation dimensions.

The gauge/matter defect $45 - 16 = 29$ is therefore an *a posteriori* cross-check, not the derivation path. The derivation path is: $\text{PSC} \rightarrow SO(10) \text{ selected} \rightarrow \text{seesaw index} = \dim(\mathbf{adj}) - \dim(\mathbf{spinor}) = 29$ by Lie theory. The separately derived FN exponent $29/9$ then agrees with this structural integer. That three independent derivation paths (topological, mirror-offset, GUT-rep) all converge on the same rational is the over-determination argument, and the GUT-rep path is the one that ties the number most directly to known group theory.

This interpretation is machine-checked as `seesaw_index_is_gauge_matter_defect` (zero sorry, module `MassRelations.SeesawIndex`, `ugp-lean`). The Lean proof verifies the equality $\dim(\mathbf{adj}_{SO(10)}) - \dim(\mathbf{spinor}_{SO(10)}) = 45 - 16 = 29$ at the level of exact arithmetic on group-theory integers; it does not presuppose the neutrino mass prediction, so the certification is genuinely pre-hoc relative to the CODATA comparison.

The three decompositions of (5) are thus all facets of the same underlying gauge/matter defect: the topological decomposition $(N_c^3 + s)$ and the mirror-offset decomposition $(4N_c^2 - \delta)$ are arithmetic rewritings of $\dim(\mathbf{adj}) - \dim(\mathbf{spinor})$, and the GUT-rep decomposition $(d_{45} - d_{16})$ is the direct statement of the defect.

4.4 Froggatt–Nielsen texture realising $b^{29/9}$

The structural decompositions of $29/9$ identify *what* the exponent encodes; we now identify the explicit Froggatt–Nielsen (FN) texture that *generates* it on the right-handed neutrino mass matrix, completing the bridge from structural arithmetic to a concrete UV completion.

Two-flavon ansatz. We adopt the same $U(1)_1 \times U(1)_2$ two-flavon FN framework used in the charged-lepton sector [12]. Flavon Φ_1 carries $U(1)_1$ charge $+1$ with VEV proportional to the Braid-Atlas b -value, $\langle \Phi_1 \rangle \propto b_g$; flavon Φ_2 carries $U(1)_2$ charge $+1$ with VEV $\langle \Phi_2 \rangle \propto b_g^{1/N_c^2}$ (the natural unit at which $\theta_{\text{Koide}} = 2/9 = 2/N_c^2$ contributes integrally). The right-handed neutrino with FN charges (q_1, q_2) acquires the Majorana mass scaling

$$M_R(g) \propto \langle \Phi_1 \rangle^{q_1} \langle \Phi_2 \rangle^{q_2} \propto b_g^{q_1 + q_2/N_c^2}, \quad (6)$$

and a seesaw light-neutrino mass $m_{\nu_g} \propto 1/M_R(g) \propto b_g^{-(q_1 + q_2/N_c^2)}$. The exponent $\alpha = -(q_1 + q_2/N_c^2)$ matches $-29/9$ when

$$q_1 + q_2/N_c^2 = 29/9. \quad (7)$$

Solutions and structural texture. Constraint (7) at $N_c = 3$ admits exactly four non-negative integer solutions in the FN-natural range ($0 \leq q_1 \leq 6$, $0 \leq q_2 \leq 30$):

$$(q_1, q_2) \in \{(0, 29), (1, 20), (2, 11), (3, 2)\}. \quad (8)$$

The completeness of this list (every solution in the range is in the set) is machine-checked in `fn_solutions_are_complete` (`MassRelations.NeutrinoFroggattNielsen`, zero sorry, [18]).

First-principles selection: MDL singleton-atomicity. The four solutions are equivalent at the level of the b-power-law constraint; the structural selection rule is provided by UGP axiom A3 (Compression / MDL). Define a charge q as *singleton-atomic in N_c* if it equals one of the elementary N_c -primitives $\{N_c, N_c-1, N_c+1, N_c^2, N_c^2-1, \text{strand}\}$, and a pair (q_1, q_2) as singleton-atomic if both charges are. These six values are not chosen to make $(3, 2)$ the answer: they are precisely the integers that appear as direct outputs of the N_c -chain in the charged-lepton sector — the strand count $\text{strand} = (N_c^2 - 1)/4$, the color-rank N_c , the $\text{SU}(2)/\text{SU}(3)$ rank boundary values $N_c \pm 1$, and the generation-quadratic values $N_c^2 \pm 1$ — all of which are independently licensed by the Braid Atlas and the VV mechanism *before* any reference to the neutrino sector [12, 15]. The singleton-atomic set is therefore the natural MDL vocabulary for UGP flavor charges, not a post-hoc construction.

Theorem 4.1 (FN texture uniqueness under MDL). *Among the four FN solutions of (7) at $N_c = 3$, the unique singleton-atomic texture is*

$$(q_1, q_2) = (N_c, \text{strand}) = (3, 2),$$

where $\text{strand} = (N_c^2 - 1)/4$ is the Braid-Atlas strand count.

Machine-checked: `fn_texture_3_2_is_unique_singleton_atomic, fn_structural_texture_existence_and_uniqueness` (*MassRelations.NeutrinoFroggattNielsen, zero sorry, [18]*). The proof enumerates the four solutions and verifies, by *decide*, that the alternatives $(0, 29)$, $(1, 20)$, $(2, 11)$ each contain a charge outside the singleton-atomic set, while $(3, 2)$ has both charges in it.

The alternative $(2, 11)$ admits the structural reading $q_2 = N_c^2 + \text{strand} = 9 + 2$ (compound, two atoms summed) or the cascade-derived reading $q_2 = b(\nu_{\mu R}) = 11$ (the Braid-Atlas b-value, itself a derived quantity from the GTE evolution). Both readings have higher MDL cost than the singleton-atomic (N_c, strand) texture, which is why Theorem 4.1 selects $(3, 2)$ on first-principles compression grounds rather than aesthetic preference.

$$q_1 = N_c, \quad q_2 = \text{strand} = (N_c^2 - 1)/4 = 2, \quad (9)$$

where $\text{strand} = \dim \mathfrak{su}(N_c)_{\text{adj}}/4$ is the Braid-Atlas strand count, machine-checked as `strand_count_eq_su_nc_adj_div_4`, zero sorry.

Cross-checks against the three structural decompositions. The texture (9) reproduces 29/9 identically under each of the three decompositions of §4:

$$\frac{q_1 N_c^2 + q_2}{N_c^2} = \frac{N_c^3 + \text{strand}}{N_c^2} = \frac{N_c(N_c^2 - 1)/2 + \delta}{(N_c^2 - 1)/2 + \delta - N_c} = \frac{45 - 16}{N_c^2} = \frac{29}{9}, \quad (10)$$

where the second equality uses $N_c^3 = N_c(N_c^2 - 1) + N_c$ and the third equality uses $\dim(\mathfrak{adj}_{\text{SO}(10)}) - \dim(\mathfrak{spinor}_{\text{SO}(10)}) = 29$. Each decomposition is therefore the same FN texture rewritten in a different algebraic basis; the texture and the structural decompositions are mutually consistent.

Computational artefact. The texture identification, the four-solution enumeration, and the cross-check against the three structural decompositions are recorded in `papers/01_SM/canonical_run/comp_p24_SP2_fn_texture_b29_9.py` (SHA-256 pre-committed prediction block; verdict `TEXTURE_3_2_FROM_Nc`; artefact JSON SHA-256 92d2b566...).

Status. The FN texture identification is structural and Lean-certified at the texture-selection level: the existence and uniqueness of the singleton-atomic solution (N_c , strand) is machine-checked (`MassRelations.NeutrinoFroggattNielsen`, zero `sorry`; see Theorem 4.1). The complementary lepton-sector FN setup `MassRelations.FroggattNielsen` provides the two-flavon framework within which this neutrino-sector texture sits. A complete first-principles derivation tying the Braid Atlas strand structure to the $\text{SO}(10)$ **126** Yukawa operator at the Lagrangian level remains open beyond the texture-selection layer; the b-power-law statement and its structural origin are unchanged from §4.

5 Absolute Mass Scale

5.1 The structural Dirac scale

The ratio $\Delta m_{21}^2/\Delta m_{31}^2$ is predicted without any absolute mass scale input (the ratio cancels the overall normalisation). An absolute scale requires identifying the Dirac Yukawa coupling magnitude.

We propose the structural Dirac scale

$$E_D = \frac{v_H}{4N_c^2 - \delta} = \frac{v_H}{29} \approx \frac{246.22 \text{ GeV}}{29} \approx 8.49 \text{ GeV}, \quad (11)$$

where $v_H = 246.22 \text{ GeV}$ is the Standard Model Higgs vacuum expectation value. The denominator 29 is the *same structural integer* that appears as the numerator of the seesaw exponent in (4). This is not a coincidence: it is Lean-certified that $29 = 4N_c^2 - \delta$ in both contexts.

5.2 Predicted mass sum

With $E_D = v_H/29$ and the Type-I seesaw formula $m_{\nu_g} = E_D^2 \cdot b_g^{29/9}/M_R$, calibrated to reproduce the observed Δm_{31}^2 at $M_R = 2 \times 10^{16} \text{ GeV}$ (within the natural GUT-scale range), we find:

$$\sum_g m_{\nu_g} \approx 56 \text{ meV} \quad (M_R = 2 \times 10^{16} \text{ GeV}), \quad (12)$$

within the Planck 2018 bound of $\sum m_\nu < 120 \text{ meV}$ [3]. The sensitivity to M_R is summarised in Table 2.

Table 2: $\sum m_\nu$ sensitivity to the GUT scale M_R with the structural Dirac scale $E_D = v_H/29$.

M_R (GeV)	$\sum m_\nu$ (meV)	Planck bound
10^{16}	≈ 102	Near limit
2×10^{16}	≈ 56	Within
3×10^{16}	≈ 40	Within
5×10^{16}	≈ 26	Within

The preferred structural benchmark is $\sum m_\nu \approx 56 \text{ meV}$ at $M_R = 2 \times 10^{16} \text{ GeV}$ (the benchmark GUT scale), with the abstract’s [55, 75] meV band set by the structural Dirac scale $E_D = v_H/29$ before any M_R calibration. Over the Planck-consistent GUT range

$M_R \in [10^{16}, 3 \times 10^{16}]$ GeV the prediction gives $\sum m_\nu \in [40, 102]$ meV, consistent with all current cosmological bounds (see Table 2). At the upper edge of the broader range $M_R = 5 \times 10^{16}$ GeV the seesaw prediction falls to ≈ 26 meV; this is within the Planck bound but would require near-maximal GUT-threshold running and sits outside the preferred $[40, 102]$ meV window, so we treat it as a sensitivity floor rather than a preferred range endpoint.

6 Individual Mass Predictions

The individual neutrino masses depend on both M_R and E_D . With the structural scale $E_D = v_H/29$ and $M_R = 2 \times 10^{16}$ GeV, anchored to reproduce Δm_{31}^2 exactly:

$$\begin{aligned} m_1 &\approx 1.5 \times 10^{-4} \text{ eV}, \\ m_2 &\approx 8.6 \times 10^{-3} \text{ eV}, \\ m_3 &\approx 5.0 \times 10^{-2} \text{ eV}. \end{aligned} \tag{13}$$

The effective Majorana mass for neutrinoless double-beta decay is $m_{\beta\beta} \approx |m_1 U_{e1}^2 + m_2 U_{e2}^2 + m_3 U_{e3}^2|$. For normal ordering with the PMNS angles from the NuFIT fit, this gives $m_{\beta\beta} \lesssim 4.5$ meV, well below the current experimental sensitivity of LEGEND-200 [1] and within reach of LEGEND-1000.

7 Preregistered Falsification Tests

The following predictions are derived with zero free parameters (ratio test) or with M_R in the natural GUT range (absolute scale tests). All are preregistered here before comparison to future experimental results.

Falsification Criteria

F1. Mass-squared ratio (parameter-free). $\Delta m_{21}^2 / \Delta m_{31}^2 = 0.02936$. A precision determination outside $[0.028, 0.031]$ (covering $\pm 3\%$ around the prediction) by JUNO, DUNE, or Hyper-Kamiokande would falsify the framework. JUNO is expected to determine the ratio to $\lesssim 1\%$ within the next decade [5].

F2. Normal mass ordering (automatic). The b -value ordering $b_1 < b_2 < b_3$ forces $m_1 < m_2 < m_3$. A confident inverted-ordering detection by DUNE or Hyper-Kamiokande would falsify the framework.

F3. Total mass sum (structural scale $E_D = v_H/29$). $\sum m_\nu \in [40, 102]$ meV for $M_R \in [10^{16}, 3 \times 10^{16}]$ GeV (the preferred Planck-consistent GUT range); the structural benchmark is ≈ 56 meV at $M_R = 2 \times 10^{16}$ GeV. A CMB-S4 or Euclid determination of $\sum m_\nu > 110$ meV or < 30 meV would falsify the structural scale identification.

F4. Effective Majorana mass (normal ordering). $m_{\beta\beta} \lesssim 5$ meV for the individual masses in (13). Any claim $m_{\beta\beta} > 10$ meV would indicate inconsistency. LEGEND-1000 will probe to ~ 10 meV [2].

F5. b -value triple uniqueness. The Braid Atlas uniquely assigns $\{5, 11, 19\}$ to the right-handed neutrinos from the GTE integer structure. Any alternative Braid

Atlas assignment producing a different $\Delta m_{21}^2/\Delta m_{31}^2$ prediction must explain why $\{5, 11, 19\}$ was not the correct assignment.

8 Discussion

8.1 What is and is not claimed

We claim:

- A *structural prediction* of $\Delta m_{21}^2/\Delta m_{31}^2 = 0.02936$ to 0.16σ from NuFIT 6.0 (0.52%) from the Braid Atlas b -values, with no free parameters. Since this is a *ratio* of mass-squared differences, it is exactly preserved under any uniform renormalization factor $Z(\mu, \mu_0)$ applied to all neutrino masses—machine-checked as `mass_ratio_Z_independent` (`MassRelations.ScaleTransport`, zero sorry). The prediction holds at any renormalization scale, not only at the oscillation-measurement scale.
- A *structural identification* of the seesaw exponent $29/9$ via three independent decompositions in N_c , δ , and GUT representation dimensions.
- A *structural scale* $E_D = v_H/29$ for the Dirac Yukawa coupling, with the denominator 29 identified by the same algebra that gives the exponent numerator.
- A *cross-sector bridge* via $\dim(126) = 2N_c^2\delta$ connecting the neutrino Majorana sector to the down-quark VV coefficient and the charged-lepton mirror offset.
- A *Froggatt–Nielsen texture identification* (§4.4): a two-flavon $U(1)_1 \times U(1)_2$ texture on the right-handed neutrino with charges $(q_1, q_2) = (N_c, \text{strand}) = (3, 2)$, both individually structural in N_c , reproduces the $b^{29/9}$ exponent. Pre-committed artefact `comp_p24_SP2_fn_texture_b29_9.json`.

We do *not* claim:

- A complete Lean-formalised derivation of the texture-to-exponent mapping from the $SO(10)$ Yukawa structure. The FN texture is identified in §4.4; an exhaustive Lean module formalising the mapping from (q_1, q_2) to the b-power-law exponent (extending the existing `MassRelations.FroggattNielsen` lepton-sector setup to the neutrino sector) is the natural next module.
- A UGP-internal determination of M_R . The GUT scale $M_R \sim 2 \times 10^{16}$ GeV is an $O(1)$ input consistent with standard gauge-coupling unification; within the UGP, M_R does not yet arise from a Lean-certified first-principles argument.

8.2 Comparison to conventional approaches

In conventional seesaw approaches, the ratio $\Delta m_{21}^2/\Delta m_{31}^2$ is entirely determined by the structure of the Dirac and Majorana mass matrices, which are not predicted without additional model-building input (texture assumptions, flavour symmetries, etc.). The UGP approach predicts R directly from the Braid Atlas topological invariants, without imposing a flavour symmetry by hand. The prediction is falsifiable at $\sim 1\%$ precision by JUNO.

8.3 Open questions

The following questions remain open:

1. **Lean formalisation of the b-power-law derivation.** The arithmetic/algebraic chain from FN texture to the mass ratio is *complete* in `ugp-lean` (`MassRelations.NeutrinoMassRatio` 2026-05-16, 5 theorems, zero sorry): FN texture $(q_1, q_2) = (3, 2)$ gives exponent $29/9$ (`fn_texture_gives_seesaw_exponent`), M_R cancels algebraically (`seesaw_ratio_independent_of_MR`), the coarse bound $0.029 < R < 0.030$ is certified by integer 9th-power comparisons (`neutrino_mass_ratio_coarse_bound`), the tight bound $|R - 0.02936| < 0.0001$ is proved via unit-width integer bounds on $b^{58/9}$ (`neutrino_mass_ratio_tight_bound`), and R is certified within 1% of the NuFIT 6.0 central value (`neutrino_mass_ratio_within_1pct_of_nufit`). What remains open is the *full Lagrangian bridge*: a formal derivation that the Yukawa couplings scale as b^{q_1} and b^{q_2} due to the FN texture, tying the Braid Atlas strand structure to the $\text{SO}(10)$ **126** Yukawa operator at the Lagrangian level. This requires extending the lepton-sector `MassRelations.FroggattNielsen` log-Yukawa formalism to the Majorana sector and will be addressed in a companion `NeutrinoLagrangian.lean` module.
2. **M_R from UGP.** Identify a UGP-internal mechanism that fixes M_R at the GUT scale without an external input.
3. **CP violation.** The Dirac CP phase δ_{CP} is not yet predicted from the Braid Atlas alone.
4. **Sterile neutrinos.** Whether the GTE integer structure admits a non-SM sterile neutrino with a distinct Braid Atlas signature is unexplored.

9 Conclusion

We have presented a parameter-free structural prediction of the neutrino mass-squared splitting ratio from the Universal Generative Principle’s Braid Atlas. The right-handed neutrino b -values $\{5, 11, 19\}$ satisfy $m_{\nu_g} \propto b_g^{29/9}$, giving $\Delta m_{21}^2 / \Delta m_{31}^2 = 0.02936$, matching NuFIT 6.0 [8] to 0.16σ (0.52%) with no free parameters and with normal ordering automatic.

The exponent $29/9$ is over-determined by three independent structural decompositions in terms of the QCD colour rank $N_c = 3$, the charged-lepton mirror offset $\delta = 7$, and the $\text{SO}(10)/\text{SU}(5)$ GUT representation dimensions — all Lean-certified in `ugp-lean` [17, 18]. The same integer 29 controls both the seesaw exponent and the structural Dirac Yukawa scale $E_D = v_H/29$, placing $\sum m_\nu \in [40, 100]$ meV within the Planck bound.

The framework makes five preregistered, experimentally falsifiable predictions (F1–F5), testable by JUNO (ratio to 1%), DUNE and Hyper-Kamiokande (ordering), CMB-S4 and Euclid (mass sum), and LEGEND-1000 (Majorana mass). The most important near-term test is the JUNO determination of $\Delta m_{21}^2 / \Delta m_{31}^2$: a result outside $[0.028, 0.031]$ would falsify the framework, while confirmation at 0.4% precision would constitute strong support.

Code and Data Availability

All computational verification scripts are available at <https://github.com/novaspihack/ugp-physics> under `papers/01_SM/canonical_run/`. The relevant scripts are: `comp_p01_EBF_17` (neutrino survey), `EBF_18` (126 bridge), `EBF_19` (29/9 derivation), `EBF_20` (absolute scale), `EBF_21` (structural decomposition), `EBF_22` (full mechanism), `EBF_24` (SO(10) CG / Majorana), and the Froggatt–Nielsen texture identification `comp_p24_SP2_fn_texture_b29_9` (SHA-256 pre-committed prediction block; verdict `TEXTURE_3_2_FROM_Nc`). Each script emits a JSON artifact with a SHA-256 hash; full filenames and descriptions are in `papers/01_SM/PROVENANCE.md` and `papers/22_neutrino_masses/PROVENANCE.md`. Lean formalization: `ugp-lean` [17, 18] (see companion formalization paper [17], `zero sorry; lake build` to verify; Lean 4.29.0-rc6, Mathlib 4.29.0-rc6). Neutrino oscillation data: NuFIT 6.0 [8] (Esteban et al., <http://www.nu-fit.org/>); NuFIT 5.2 [9] (historical comparison).

Acknowledgments

AI assistants were used for adversarial review, coding assistance, and debugging during manuscript preparation and code development. All scientific claims, derivations, and numerical results were independently verified by deterministic scripts and Lean proofs; no AI-generated content appears as unverified scientific output.

References

- [1] N. Abgrall et al. The LEGEND-200 and LEGEND-1000 experiments. *J. Phys. Conf. Ser.*, 1468:012136, 2020. arXiv eprint unverified; the arXiv for LEGEND-1000 design report is 2107.11462 (distinct paper). Check INSPIRE for the correct eprint for this proceedings contribution.
- [2] N. Abgrall et al. LEGEND-1000 preconceptual design report. 2021.
- [3] N. Aghanim et al. Planck 2018 results. VI. Cosmological parameters. *Astron. Astrophys.*, 641:A6, 2020.
- [4] M. Aker et al. Direct neutrino-mass measurement based on 259 days of KATRIN data. *Science*, 388:adq9592, 2025. Published April 11, 2025 (arXiv preprint June 2024; key retains 2024 for backward compatibility).
- [5] F. An et al. Neutrino physics with JUNO. *J. Phys. G*, 43:030401, 2016.
- [6] Sundance O. Bilson-Thompson. A topological model of composite preons. *arXiv preprint*, 2005. arXiv submitted March 22, 2005 (originally listed year 2006 in bib was wrong).
- [7] Sundance O. Bilson-Thompson, Fotini Markopoulou, and Lee Smolin. Quantum gravity and the standard model. *Classical and Quantum Gravity*, 24(16):3975–3993, 2007.

- [8] I. Esteban, M.C. Gonzalez-Garcia, M. Maltoni, I. Martinez-Soler, J.P. Pinheiro, and T. Schwetz. NuFit-6.0: updated global analysis of three-flavor neutrino oscillations. *JHEP*, 12:216, 2024. NuFIT 6.0 global fit (2024), <http://www.nu-fit.org/>.
- [9] I. Esteban, M.C. Gonzalez-Garcia, M. Maltoni, T. Schwetz, and A. Zhou. NuFIT 5.2 (2022). *JHEP*, 09:178, 2020. Updated results available at <http://www.nu-fit.org/>.
- [10] M. Gell-Mann, P. Ramond, and R. Slansky. Complex spinors and unified theories. In P. van Nieuwenhuizen and D. Freedman, editors, *Supergravity*, pages 315–321. North-Holland, 1979.
- [11] P. Minkowski. $\mu \rightarrow e\gamma$ at a rate of one out of 10^9 muon decays? *Phys. Lett. B*, 67:421–428, 1977.
- [12] N. Spivack. Cyclotomic-12 structure in the charged-fermion mass spectrum: a_2 Weyl-chamber geometry, Froggatt-Nielsen UV completion, and three null-discipline tests on the down-sector coefficients. Draft manuscript; source at <https://github.com/novaspivack/ugp-physics>, 2026. Zenodo DOI placeholder — to be updated before submission. Source: <https://github.com/novaspivack/ugp-physics>.
- [13] N. Spivack. A deterministic number-theoretic framework for the standard model parameter spectrum, 2026. Zenodo DOI placeholder — to be updated before submission. Source: <https://github.com/novaspivack/ugp-physics>.
- [14] N. Spivack. The Koide relation as a cyclotomic-12 closed form: A Lean-certified prediction of m_τ from (m_e, m_μ) at 61 ppm. Zenodo, 2026. Zenodo DOI placeholder — to be updated before submission.
- [15] Nova Spivack. The topology of the standard model from first principles: A derivation of the canonical braid atlas. Zenodo, 2025.
- [16] Nova Spivack. PSC-optimality of the standard model in 4D gauge theory space. NEMS/PSC Suite, Paper 05, 2026. <https://doi.org/10.5281/zenodo.19429721>. Two-Layer PSC Theorem: hard constraints force $SU(3) \times SU(2) \times U(1)$; Presentation Invariance selects 3 generations. Machine-checked in nems-lean.
- [17] Nova Spivack. ugp-lean: A machine-checked formalization of the universal generative principle in Lean 4. Zenodo, 2026. Formalization paper (latest version): <https://doi.org/10.5281/zenodo.19554688>. Concept DOI (always-latest): <https://doi.org/10.5281/zenodo.19433538>. 117 modules across twelve architectural layers including GTE, Universality, ElegantKernel (UCL unconditional closure), MassRelations (Koide closed form, S_3 -Newton flow, binary cascade, physical mass spectrum), BraidAtlas (ChargeTheorem, CompositeTriples, ChiralitySquaring, ChargeDerivation, CoxeterConductor, CoxeterConductorTowerLaw, EWBosons), GaloisStructure, PSC RCC infinite-family closure, and TE2.2 scan certification. Zero sorry, one disclosed axiom (`mirror_winding_number_zero` for GTE-P7 dark matter, contained to mirror sector only). Lean 4.29.0-rc6 / Mathlib 4.29.0-rc6.
- [18] Nova Spivack. ugp-lean: Lean 4 formalization of UGP and GTE. <https://github.com/novaspivack/ugp-lean>, 2026. 117-module Lean 4 library (module count updated as of 2026-05-09). lake build to verify (zero sorry, zero custom

axioms, standard Mathlib axiom signature). Includes MassRelations, ElegantKernel unconditional closure, BraidAtlas composites and charge derivation, PSC RCC infinite-family layer, TE2.2 scan certification. New in GTE.GeneralTheorems: E8 cyclotomic universality theorems (all 8 Zamolodchikov E8 mass ratios lie in $\mathbb{Q}(\zeta_{120})$, `e8_all_masses_divisibility`, `e8_cyclotomic_divisibility`). Lean 4.29.0-rc6 / Mathlib 4.29.0-rc6.

- [19] T. Yanagida. Horizontal gauge symmetry and masses of neutrinos. In *Proceedings of the Workshop on the Baryon Number of the Universe and Unified Theories*, pages 95–99, 1979.