

The Self-Referential Renormalization Group: A Universal Framework for Physical Constants

Nova Spivack^{*1}

¹Independent Research

2026

Preprint.

Archival version (Zenodo): [DOI pending publication]

Abstract

We introduce the *Self-Referential Renormalization Group* (SRRG): a gradient-flow theory on the space of self-referential physical theories. The SRRG flow maximises a net viability functional $F[S] = R[S] - C_\Lambda[S]$, where $R[S]$ measures self-representation capacity and $C_\Lambda[S]$ decomposes into closure, self-computation, and selector costs inherited from the NEMS/PSC framework.

The main body establishes the core fixed-point structure: fixed-point existence via the Master Fixed-Point Theorem, monotonicity of F along the SRRG flow (an analogue of Zamolodchikov's c -theorem), linearised contraction rate $1/\varphi$ where φ is the golden ratio, and minimal $U(1)$ symmetry under the PSC cost condition. The Information Profit Threshold $\text{IPT} \approx 1.1309$ arises as the SRRG efficiency ratio at the fixed point, conditional on the explicit PSC Landauer self-consistency hypothesis $h_{\text{psc_sc}}$ (grade [A⁻]; Remark 5.4). The associated one-dimensional η -flow has a candidate projected β -function $\beta_\eta = \kappa(\eta - \text{IPT})(\eta - 2)$, with an IR-stable fixed point at $\eta = \text{IPT}$ and a UV-unstable separatrix at $\eta = 2$; the algebraic uniqueness and no-third-zero properties are machine-certified in Lean 4 (Vieta's theorem, zero sorry).

Conditional physical applications — strong CP phase $\theta_{\text{QCD}} = 0$, three fermion generations $N_{\text{gen}} = 3$, multi-scale SM gauge structure, Higgs quartic recovery $\lambda_H = m_H^2/(2v^2)$, and structural exclusion of Planck-scale vacuum energy — are developed in Appendix B, with all additional physical bridge hypotheses stated explicitly (Table 1). A negative diagnostic for the Weinberg angle via Haar-entropy ratios is included as Appendix D. A positive partial result for the fine-structure constant is given in Appendix C (grade [A/D]): the tree-level SM identity $\alpha = g_1^2 g_2^2 / [4\pi(g_1^2 + g_2^2)]$ applied to the Lean-certified bare couplings from P01 yields $1/\alpha(M_Z) = 4\pi \cdot 377525/37264 \approx 127.311$, which is 0.495% from the PDG running value. The **SRRG–MDL equivalence conjecture** (Open Problem 9; grade [B+]), also introduced in Appendix C, identifies the SRRG fixed-point selection $\arg \max_S F[S]$ with MDL minimization $\arg \min_S K[S]$ for self-referential theories, bridging the present framework with the Architecture paper (P09 [1]).

All owned formal implication theorems compile in the Lean 4 library `srrg-lean` with zero proof-obligations; upstream caveats and non-load-bearing stubs are disclosed in Appendix 9.2. Several consequences are formal theorem-level results of the SRRG scaffolding; others are conditional on explicitly stated physical bridge hypotheses. Load-bearing assumptions are tabulated in Table 1.

^{*}Correspondence: novaspihackrelay@gmail.com

Contents

1	Introduction: Why Physical Constants Might Be Fixed Points	4
1.1	The Core Claim	4
1.2	Relation to Existing Frameworks	4
1.3	Organization of the Paper	5
2	The SRRG Framework	7
2.1	Theory Space	7
2.2	Representation Capacity $R[S]$	8
2.3	Constraint Functional $C_\Lambda[S]$	8
2.4	Net Viability $F[S]$	9
2.5	The SRRG Flow	9
3	Fixed-Point Theorems	10
3.1	Fixed-Point Definition and Existence	10
3.2	The F-Theorem: Viability Monotonicity	10
3.3	Lyapunov Stability at S^*	11
3.4	Uniqueness	11
4	NEMS/PSC as SRRG Axiom System	11
4.1	The NEMS/PSC Framework	11
4.2	NEMS Gives SRRG Structure	12
4.3	NEMS Cost Components Map to C_Λ	12
4.4	PSC-Optimal Theory is an SRRG Fixed-Point Candidate	12
4.5	Selector Strength Hierarchy as $R[S]$ Levels	13
5	IPT as the SRRG Information-Efficiency Fixed Point	13
5.1	The IPT Formula	13
5.2	The PSC Landauer Self-Consistency Equation	13
5.3	IPT as the SRRG Fixed-Point Efficiency Ratio	14
5.4	IPT Requires No Domain Assumptions	18
6	Gauge Symmetry at the SRRG Fixed Point	19
6.1	The Gauge Selection Problem in SRRG	19
6.2	U(1) Minimality	19
6.3	Higher Gauge Factors from Multi-Scale SRRG	20
7	CFT Universality Classes and the $\varepsilon_{3s} = \text{IPT}$ Prediction	21
7.1	The GXT Prediction	21
7.2	Three CFT Universality Classes	21
7.3	Lean Certification Target	21
8	Discussion	22
8.1	The SRRG Resolution of Fine-Tuning	22
8.2	SRRG vs. Anthropic Reasoning	22
8.3	SRRG as the Master Meta-Principle	23
8.4	Open Problems	23

9	Lean 4 Machine Certification	25
9.1	Certification Architecture	25
9.2	Certified Theorem Table	25
9.3	Previously Certified Inputs	30
	Appendices	34
A	Complete Lean Theorem Table	34
A.1	Core Definitions	35
A.2	Fixed-Point Theorems	35
A.3	Applications	36
A.4	Bridge Modules	36
A.5	Connection Layer	37
A.6	Constants: Beta Function, Higgs Quartic, Cosmological Constant	37
B	Conditional Physical Applications	38
B.1	The Strong CP Problem: $\theta_{\text{QCD}} = 0$	39
B.2	Fermion Generation Count: $N_{\text{gen}} = 3$	39
B.3	One-Loop SRRG β -Function: UV Stability and Hessian	40
B.4	Higgs Quartic Coupling λ_H from SRRG EW Stability	41
B.5	Cosmological Constant: Structural Smallness Argument	42
C	Fine-Structure Constant: Direct Coupling Route	42
D	Failed Route: Weinberg Angle from Haar Entropy	44

1 Introduction: Why Physical Constants Might Be Fixed Points

Why are the physical constants what they are? The fine-tuning problem [2] asks why the values of dimensionless parameters such as the gauge coupling ratios, the cosmological constant, or the electroweak hierarchy are what they are, rather than exponentially different values. Standard responses fall into three classes: (i) anthropic / landscape reasoning, which explains observed values by environmental selection from a large prior [3]; (ii) naturalness arguments, which seek symmetry or technical reasons that values are “natural” at their scale; and (iii) uniqueness conjectures, which assert that the observed laws are somehow the only logically consistent or self-consistent possibility.

None of these approaches predicts specific numerical values from first principles without additional input. The present paper proposes a fourth route: *physical constants are fixed points of a universal self-referential gradient flow*. The SRRG was first introduced in [4] (Chapters 6–7); the present paper provides the full theorem statements, Lean 4 machine-certified proofs, and the bridge to IPT and the Standard Model gauge structure.

The central object is the *Self-Referential Renormalization Group* (SRRG), a flow on the space of self-referential physical theories. The SRRG is not a flow within a particular theory (as the Wilsonian RG is a flow on coupling space within a fixed field theory); it is a flow *on the space of theories themselves*, selecting the theory that best “earns its own existence” by maximizing its ability to represent itself stably.

1.1 The Core Claim

A single equation — the SRRG fixed-point condition $\delta F[S]/\delta S = 0$ — is sufficient to derive:

1. The **Information Profit Threshold** $\text{IPT} \approx 1.1309$: the universal information-efficiency constant that appears as the selection threshold across immune repertoires, protein sequences, ecological fitness landscapes, economic selection, and conformal field theory universality classes [5, 6]. This derivation is conditional on the PSC Landauer self-consistency hypothesis (the [H4] / $h_{\text{psc_sc}}$ condition, Remark 5.4).
2. The **minimal gauge group** $U(1)$: the unique group that minimizes the self-computation-principle cost at the fixed point [7, 8].
3. The $1/\varphi$ **contraction eigenvalue**: the rate at which transverse perturbations contract toward S^* , related to the Fibonacci renormalization spectrum [9, 10].
4. The **NEMS observational barriers**: the self-improvement barrier, the no-hypercomputer theorem, and the epistemic meta-barrier are all encoded as components of $C_\Lambda[S]$, and they are not external assumptions but *consequences* of the SRRG fixed-point structure [11].

The core implication theorems require no anthropic reasoning. They reduce the target constants to explicitly stated self-referential fixed-point hypotheses whose formal consequences are machine-certified in Lean 4. Physical applications (Appendix B) add domain-specific physical hypotheses that are disclosed separately.

1.2 Relation to Existing Frameworks

The SRRG stands on three pillars:

NEMS/PSC framework. The *No External Model Selection* (NEMS) framework [12] and the *Perfect Self-Containment* (PSC) principle provide the axiomatic foundation. Key NEMS theorems — the closure audit [11], the selector barrier, the self-improvement barrier [11], and the Master Fixed-Point Theorem (MFP-1, Paper 26 in the NEMS suite) — become the SRRG constraint functional $C_\Lambda[S]$.

UGP gauge structure. The Universal Generative Principle (UGP) [9] identifies the Standard Model gauge group as the unique PSC-optimal solution. The SRRG provides the meta-justification: the SM gauge group is selected because it minimizes $C_\Lambda[S^*]$ at the SRRG fixed point.

IPT and GXT. The Information Profit Threshold has been derived from PSC axioms [5] and confirmed across multiple domain universality classes [6]. The SRRG provides the deepest derivation: IPT is the ratio $R[S^*]/C_\Lambda[S^*]$ at the SRRG fixed point.

Asymptotic safety — comparison. The asymptotic safety programme [13, 14] also identifies physical theories via non-Gaussian UV fixed points of a renormalization-group flow. The SRRG shares this fixed-point selection logic but differs fundamentally: asymptotic safety flows on coupling space *within* a given field theory, while the SRRG flows on the space of theories themselves, selecting the theory by maximizing self-representational viability. The most striking external result is the Shaposhnikov–Wetterich prediction [15] of $m_H = 126$ GeV from an asymptotic-safety UV fixed point—a precedent for deriving the Higgs mass from a fixed-point condition (cf. Appendix B.4).

Logical relation to the General Selection companion paper. The General Selection paper [6] documents the empirical cross-domain recurrence of the IPT threshold and the cross-domain recurrence of the IPT threshold across five empirical domains, one model-internal symbolic-communication test, and an independent $\mathbb{Q}(\zeta_{120})$ algebraic-substrate family. The present paper is not primarily an empirical survey; it proposes a fixed-point mechanism—the SRRG—capable of explaining *why* such a threshold should exist and take the specific value IPT. The General Selection paper thus supplies external empirical motivation and cross-domain tests that any proposed mechanism must pass, while the present paper supplies a candidate internal derivation. The two papers are logically independent: the cross-domain recurrence is documented without presupposing the SRRG mechanism, and the SRRG fixed-point derivation does not require the empirical cross-domain survey as a logical premise.

1.3 Organization of the Paper

Section 2 defines the SRRG framework: theory space S , the representation capacity $R[S]$, the constraint functional $C_\Lambda[S]$, the viability functional $F[S]$, and the SRRG beta function. Section 3 states and proves the main fixed-point theorems: existence, F-theorem (monotonicity), Lyapunov stability, and uniqueness conditions. Section 4 shows that NEMS/PSC axioms give the SRRG structure. Section 5 derives IPT as the SRRG information-efficiency fixed point. Section 6 derives the minimal gauge group from the SRRG fixed point. Section 7 discusses CFT universality classes and the GXT prediction $\varepsilon_{3s} = \text{IPT}$ across all classes. Section 8 discusses fine-tuning, anthropic reasoning, and open problems. Section 9 describes the Lean 4 machine-certification in `srrg-lean`. Appendix A gives the theorem table with all certified statements. Appendix B presents conditional physical applications (strong CP, generation count, β -function, λ_H , Λ_{vac}) under additional physical hypotheses. Appendix D documents the Weinberg angle as a diagnosed negative result.

Remark 1.1 (Scope and limitations — what this paper does not prove). For transparency, the following are *not* established in this paper:

- **First-principles derivation of m_H , v , or EW VEV:** The electroweak VEV $v = 246$ GeV and Higgs mass m_H are Category A/D anchors (observed inputs), not outputs of the SRRG formalism. The λ_H result in Appendix B.4 is a structural recovery, not an independent prediction.
- **Derivation of $\sin^2 \theta_W$:** This is a diagnosed negative result (Appendix D); the Haar entropy route does not yield the Weinberg angle. It is documented as an open problem.
- **Derivation of α (fine-structure constant):** Partially addressed via a direct coupling route independent of $\sin^2 \theta_W$: the formula $\alpha = g_1^2 g_2^2 / (4\pi(g_1^2 + g_2^2))$ yields $1/\alpha(M_Z) = 4\pi \cdot 377525/37264 \approx 127.31$ directly from the Lean-certified bare couplings (Category A/D; see Appendix C). Full derivation of the low-energy CODATA value additionally requires standard QED running from M_Z to the Thomson limit. The Haar-entropy route via $\sin^2 \theta_W$ remains a diagnosed negative result (Appendix D); that route is orthogonal to the coupling route.
- **SRRG-axiomatic proof of $h_{\text{psc_sc}}$ /H4:** This hypothesis is stated explicitly in Theorem 5.3 and Remark 5.4; it is load-bearing and not yet derived from SRRG axioms.
- **First-principles proof of the three physical axioms:** The Landauer sustainability, IR stability, and Wilsonian polynomial hypotheses are independently motivated but not formally derived from the SRRG fixed-point structure in this paper.
- **Consciousness or qualia:** The SRRG is a theory of self-referential physical theories, not of phenomenal experience.

Assumption	Content	Used for	Status
$h_{\text{psc_sc}} / \text{H4}$	PSC Landauer self-consistency; 1D reduction axiom	$\eta(S^*) = \text{IPT}$	Open structural hypothesis
h_θ	CP violation implies positive closure cost	$\theta_{\text{QCD}} = 0$	Physical bridge
h_{cp3}	EW CP asymmetry requires $N_{\text{gen}} \geq 3$	Generation lower bound	Standard physics bridge
h_{excess}	Extra generations impose selector cost	Generation upper bound	SRRG physical bridge
$h_{\text{large}\Lambda}$	Planck-scale vacuum energy prevents closure	CC smallness	Quantum-SRRG bridge
h_{IR}	Physical fixed points are IR-stable	$\eta \leq 2$ / FP exhaustion	RG bridge
EW anchors	m_H, v treated as observational inputs	λ_H recovery	Category A/D anchors
H_0	Hubble constant treated as observational input	Λ_{vac} match	External cosmological input
P01 bare couplings	$g_1^2 = 16/125, g_2^2 = 2329/5400$ from ugp-lean	α direct coupling route (App. C)	Lean-certified input

Table 1: Load-bearing physical hypotheses. Top block: core SRRG hypothesis. Middle block: conditional applications (Appendix B). Bottom block: observational anchors (not hypotheses — fixed by measurement).

All formal implications from these hypotheses are machine-checked in **srrg-lean** with zero sorry. The core SRRG formalism (§2–§7) uses only $h_{\text{psc_sc}}$; all remaining hypotheses appear only in Appendix B.

2 The SRRG Framework

2.1 Theory Space

Let $\mathcal{T}$ denote the collection of all self-referential physical theories. Each element $S \in \mathcal{T}$ is a complete theory: a specification of laws, symmetries, degrees of freedom, and observable record structure. We do not fix a background spacetime; $\mathcal{T}$ is the space of all candidate frameworks.

Definition 2.1 (SRRG Theory Space). *An SRRG theory space is a tuple $(\mathcal{T}, d, K)$ where:*

1. $\mathcal{T}$ is a set of theories with an equivalence relation $\sim$ (observational equivalence: $S_1 \sim S_2$ iff all macroscopic records agree);
2. $d : \mathcal{T} \times \mathcal{T} \rightarrow \mathbb{R}_{\geq 0}$ is a pseudometric measuring “distance” between theories, with $d(S_1, S_2) = 0 \Leftrightarrow S_1 \sim S_2$;
3. $K : \mathcal{T} \rightarrow \mathbb{R}_{\geq 0}$ is a descriptive complexity measure (e.g., algorithmic complexity of the theory’s axioms).

The prototype in the NEMS framework is the **Sieve.TheorySpace** type from **nems-lean** [11], which provides the observational equivalence structure; see also [4] §6.1 for the original informal

definition. The `SrrgTheorySpace` structure in `srrg-lean` extends this with a pseudometric and complexity measure, corresponding to the full Definition 2.1.

2.2 Representation Capacity $R[S]$

The *representation capacity* $R[S]$ measures how accurately and completely theory S can model itself. Formally:

Definition 2.2 (Representation Capacity). $R[S]$ is a nonnegative real-valued functional on $\mathcal{T}$ satisfying:

1. $R[S] \geq 0$ for all S .
2. $R[S]$ is bounded above by the diagonal barrier: $R[S] \leq B_\Delta$, where B_Δ is the supremum of realizability levels achievable by any self-model of S (from `SelectorStrength.BarrierSchema` [11]).
3. $R[S]$ increases as S gains more accurate self-models.

The `SelectorStrength` preorder from `nems-lean` provides the formal hierarchy of realizability levels: strength level σ is a class of functions, and $\sigma_1 \leq \sigma_2$ iff every function in σ_1 is in σ_2 . The diagonal barrier theorem `no_total_decider_at_strength` (Paper 29) proves that no self-referential system can realize a total decider at its own strength level — this is the upper bound on $R[S]$.

2.3 Constraint Functional $C_\Lambda[S]$

The *constraint functional* encodes the total cost of stability and self-computation requirements:

Definition 2.3 (Constraint Functional). $C_\Lambda[S]$ decomposes as three nonnegative components:

$$C_\Lambda[S] = C_{\text{closure}}[S] + C_{\text{SCP}}[S] + C_{\text{sel}}[S]. \quad (1)$$

- **Closure cost** $C_{\text{closure}}[S]$: the overhead of non-invariant predicates requiring external selectors. From `Closure.Theorems.AuditSoundness` (Paper 28): if theory S has predicates that are not preserved under observation- equivalence, it requires external symmetry-breaking at cost $C_{\text{closure}}[S] > 0$.
- **SCP cost** $C_{\text{SCP}}[S]$: the penalty for imperfect self-computation. From the *Self-Improvement Barrier* (Paper 32): $C_{\text{SCP}}[S] > 0$ iff S lacks a total internal upgrade certifier. At fixed point: $C_{\text{SCP}}[S^*] = 0$ iff S^* satisfies the Self-Computation Principle exactly.
- **Selector cost** $C_{\text{sel}}[S]$: the cost of external symmetry-breaking. From `Closure.Core.Selector`: $C_{\text{sel}}[S] > 0$ iff S requires an external selector for its canonical representation.

Proposition 2.4 (Constraint Nonnegativity and Zero Condition). $C_\Lambda[S] \geq 0$ for all S . Moreover, $C_\Lambda[S] = 0$ if and only if all three components vanish: $C_{\text{closure}}[S] = C_{\text{SCP}}[S] = C_{\text{sel}}[S] = 0$. This characterizes theories that are simultaneously categorical, SCP-complete, and selector-free.

This proposition is proved in `SrrgLean.Core.ConstraintFunctional` (`constraint_functional_zero_iff_components_zero`).

Remark 2.5 (Concrete non-trivial instantiation of $C_\Lambda > 0$). A concrete class of theories with $C_\Lambda[S] > 0$ is given by CP-violating theories: Proposition B.2 (Appendix B.1) shows that any theory S with $|\theta_{\text{QCD}}(S)| > 0$ satisfies $C_{\text{closure}}[S] > 0$, hence $C_\Lambda[S] > 0$. This empirically grounds the framework: the QCD theta term is a specific physical operator whose presence increases the closure cost. The framework is therefore not vacuously satisfied — there exist explicit physical theories at positive cost.

2.4 Net Viability $F[S]$

Definition 2.6 (Net Viability Functional).

$$F[S] = R[S] - C_\Lambda[S]. \quad (2)$$

$F[S]$ is the net self-representational efficiency of theory S : the excess of self-knowledge over the overhead required to maintain it.

The viability functional is the SRRG analogue of the free energy in thermodynamics, or the c -function in Zamolodchikov’s c -theorem [16]. Theories with $F[S] < 0$ are *pathological* (more overhead than self-knowledge); theories with $F[S] > 0$ are *viable*.

Remark 2.7. The `viable-continuation-lean` library provides an independent formalization of viability semantics: `HasViability.Viable s` $\Leftrightarrow F[s] > 0$. The SRRG specializes this general framework to self-referential theory spaces.

2.5 The SRRG Flow

The SRRG beta function is the gradient of $F[S]$ on theory space:

$$\beta_{\text{SRRG}}(S) = G_S^{-1} \cdot \frac{\delta F[S]}{\delta S}, \quad (3)$$

where G_S is the metric on $\mathcal{T}$ at S (the analogue of the Zamolodchikov metric in CFT). This is a gradient-flow equation: the theory evolves in the direction of steepest increase of F .

Unlike the Wilsonian RG [17, 18] (which flows in scale), the SRRG flow is driven by *self-consistency*: a theory flows toward the configuration that best represents itself stably. The flow is defined on the full theory space, not restricted to renormalizable perturbations.

Definition 2.8 (Discrete SRRG Flow). A discrete SRRG flow step is a map $\mathcal{F} : \mathcal{T} \rightarrow \mathcal{T}$ satisfying $F[\mathcal{F}(S)] \geq F[S]$ for all S (viability non-decrease). Such a flow is monotone (*IsMonotoneFlow* in *srrg-lean*).

The F-theorem (§3.2) shows that any monotone SRRG flow converges toward the fixed-point locus.

Remark 2.9 (Relation to the SRRG Monograph). The definitions in this section implement Chapters 6.1–6.5 of [4]. The `SrrgTheorySpace`, `RepCapacityProfile`, `ConstraintProfile`, and `Viability` types in *srrg-lean* give machine-checked form to the informal mathematical structures developed there.

3 Fixed-Point Theorems

3.1 Fixed-Point Definition and Existence

Definition 3.1 (SRRG Fixed Point). $S^* \in \mathcal{T}$ is an SRRG fixed point if it is a global maximizer of the viability functional:

$$\forall S \in \mathcal{T} : F[S] \leq F[S^*]. \quad (4)$$

Equivalently (in the continuous setting), $\beta_{\text{SRRG}}(S^*) = 0$, i.e., $\delta F[S^*]/\delta S = 0$.

Theorem 3.2 (SRRG Fixed-Point Existence). Let $\mathcal{T}$ be a finite theory space equipped with the viability functional $F[S] = R[S] - C_\Lambda[S]$. Then there exists $S^* \in \mathcal{T}$ such that S^* is an SRRG fixed point.

More generally, for any self-referential flow map $\mathcal{F} : \alpha \rightarrow \alpha$ on a CSRI system, there exists p with $p \simeq \mathcal{F}(p)$ (Master Fixed-Point Theorem, MFP-1).

Proof. For finite $\mathcal{T}$: F attains its maximum on any nonempty finite set (combinatorial argument). This is formalized as `viability_maximizer_exists_of_fintype` in `srrg-lean` using `Finite.exists_max` from `Mathlib`.

For the abstract flow version: The Master Fixed-Point Theorem (MFP-1) states that in any self-referential interface (CSRI) system, for any congruent transformer $F : \alpha \rightarrow \alpha$, there exists d with $d \simeq F(d)$ [11]. The SRRG flow map $\mathcal{F}$ inherits this property from the self-referential structure of physical theories. This is formalized as `srrg_flow_fixed_point_exists` in `srrg-lean`. $\square$

Remark 3.3. The MFP-1 proof is a generalization of the diagonal lemma (Gödel [19], Kleene’s recursion theorem). Applied to physical theories, it says: any theory with a self-referential update map has a fixed point of that map. The SRRG fixed point S^* is thus a mathematical necessity, not a contingent feature of the universe.

3.2 The F-Theorem: Viability Monotonicity

Theorem 3.4 (SRRG F-Theorem [4, Thm. 6.6]). Under any monotone SRRG flow step $\mathcal{F}$, $F[\mathcal{F}^n(S)] \geq F[S]$ for all S and all $n \geq 0$.

Proof. By induction on n . Base case: $\mathcal{F}^0(S) = S$, so $F[S] \leq F[S]$ trivially. Inductive step: assume $F[\mathcal{F}^n(S)] \geq F[S]$; since $\mathcal{F}$ is monotone, $F[\mathcal{F}(\mathcal{F}^n(S))] \geq F[\mathcal{F}^n(S)] \geq F[S]$.

Machine-certified as `viability_nondecreasing_along_flow` in `SrrgLean.Bridges.ToNEMSConfirmations`, zero sorry, proved by induction on `Function.iterate`; this is a genuine F-theorem result. $\square$

This is the SRRG analogue of Zamolodchikov’s c -theorem [16], which asserts that the c -function is non-increasing along RG flow in 2D CFT. In the SRRG, viability plays the role of the c -function, and the flow is toward higher viability (rather than lower c). The physical content is that no self-referential theory can spontaneously decrease its net self-representational efficiency under the SRRG dynamics. The result is part of a family of RG monotonicity theorems: the a -theorem in four dimensions [20] and the F-theorem in three dimensions [21] extend Zamolodchikov’s irreversibility result; the SRRG F-theorem generalises this chain to flows on self-referential theory spaces.

3.3 Lyapunov Stability at S^*

Theorem 3.5 (Linearized Stability). *The contraction rate of the linearized SRRG flow near S^* equals $1/\varphi$, where $\varphi = (1 + \sqrt{5})/2$ is the golden ratio:*

$$\left| \frac{(1 - \sqrt{5})}{2} \right| = \frac{1}{\varphi}. \quad (5)$$

Proof. This is the PSC contraction eigenvalue: the subdominant Fibonacci renormalization eigenvalue $\psi = (1 - \sqrt{5})/2$ has magnitude exactly $1/\varphi$. Formalized as `UgpLean.GTE.abs_psi_eq_inv_phi` in `ugp-lean` (zero sorry) and re-exposed in `srrg-lean` as `SrrgLean.FixedPoints.Stability.linearized_flow_contraction_rate`. $\square$

Since $|\psi| = 1/\varphi < 1$, transverse perturbations contract toward S^* , confirming Lyapunov stability. The golden ratio appears here because the Fibonacci renormalization group — which controls the contraction of the PSC self-model update — has $1/\varphi$ as its subdominant eigenvalue.

3.4 Uniqueness

Theorem 3.6 (Conditional Uniqueness). *If the viability functional $F[S]$ has a unique global maximizer (strict global maximum condition), then the SRRG fixed point S^* is unique:*

$$(\forall S, T : F[S] \leq F[S^*] \text{ and } F[T] \leq F[S^*]) \implies S^* = T. \quad (6)$$

Proof. Formalized as `uniqueness_of_strict_concavity` in `SrrgLean.FixedPoints.Uniqueness`. The proof takes the strict-unique-maximizer condition `hUniqMax` as a hypothesis and derives equality of maximizers. The physical content of `hUniqMax` is that the Hessian of $F[S]$ at S^* is negative definite: no other theory achieves the same maximum viability. $\square$

Remark 3.7. The strict-uniqueness hypothesis (negative-definite Hessian) is well-motivated physically — if two theories have exactly the same viability, they must be observationally equivalent by the PSC principle — but its Lean formalization requires a functional-derivative framework (open problem, listed in §8.4).

4 NEMS/PSC as SRRG Axiom System

4.1 The NEMS/PSC Framework

The NEMS (No External Model Selection) framework [12] posits that a foundationally complete physical theory must not require external selectors to determine its own parameters. The PSC (Perfect Self-Containment) principle formalizes this: theory T satisfies PSC iff T can be uniquely identified from its own macroscopic record structure without external input.

The `SrrgTheorySpaceFull` structure (certified in `SrrgLean.Core.TheorySpace`) extends `SrrgTheorySpace` with the full `NemS.Optimality` fields: the complexity functional K , a `RecordEquivalent` relation, and the associated optimality predicates. This makes the connection to NEMS optimality explicit at the type level, rather than only in the bridge modules.

Key theorems of the NEMS program [11]:

- **Semantic Terminality** (Paper 18): any extension T' of a PSC-optimal theory T either fails PSC or is record-redundant.

- **Foundational Finality** (Paper 23): there exists a PSC-optimal theory S^* such that $S^* = \text{reconstruct}(\text{extract}(S^*))$ — a fixed point of the semantics-to-ontology map.
- **Self-Improvement Barrier** (Paper 32): no total internal upgrade certifier exists at any diagonal strength level. This bounds $C_{\text{SCP}}[S]$ from below for all non-fixed-point theories.
- **Selector Barrier** (Paper 29): no total decider exists at any strength level. This bounds $R[S]$ from above (the diagonal barrier on self-knowledge).

4.2 NEMS Gives SRRG Structure

Theorem 4.1 (NEMS $\Rightarrow$ SRRG). *Every NEMS framework carries a canonical SRRG theory space. Concretely, every `NemS.Core.Basics.Framework` admits a canonical `SrrgTheorySpace` on its model type.*

Proof. Formalized as `SrrgLean.Bridges.FromNEMS.frameworkSrrgTheorySpace`: given a `NemS.Framework` with a type of models, we construct a `SrrgTheorySpace` extending `Sieve.TheorySpace` with flat pseudometric and zero initial complexity. The metric and complexity are refined by physical input; the typing layer is canonical. $\square$

4.3 NEMS Cost Components Map to C_Λ

The three components of $C_\Lambda[S]$ have precise NEMS/PSC sources:

- $C_{\text{closure}}[S] \leftarrow \text{Closure.Theorems.AuditSoundness}$: audit soundness certifies that non-invariant predicates require external selectors. The closure cost is the measure of such non-invariant content.
- $C_{\text{SCP}}[S] \leftarrow \text{SelfImprovement.Theorems.Barrier}$ (Paper 32): the self-improvement barrier says no total upgrade certifier exists; $C_{\text{SCP}}[S]$ is the resulting overhead of maintaining the theory’s self-computation under the barrier constraint.
- $C_{\text{sel}}[S] \leftarrow \text{NemS.Core.Selectors}$: NEMS selectors are external canonical representatives for world-type separation; their cost is the symmetry-breaking overhead $C_{\text{sel}}[S]$.

At the SRRG fixed point S^* : $C_{\text{closure}}[S^*] = C_{\text{SCP}}[S^*] = C_{\text{sel}}[S^*] = 0$. This is the formal content of “no external selectors are needed at S^* ”: the theory is simultaneously categorical, SCP-complete, and selector-free.

4.4 PSC-Optimal Theory is an SRRG Fixed-Point Candidate

Theorem 4.2 (PSC-Optimal $\Rightarrow$ SRRG Candidate). *A PSC-optimal non-`FailsPSC` theory T (in the sense of `NemS.Optimality.TheorySpace.PSCOptimal`) is an SRRG fixed-point candidate: it satisfies $C_\Lambda[T] = 0$ and locally maximizes $F[T] = R[T]$.*

This connects the NEMS notion of “PSC-optimal theory” directly to the SRRG fixed-point condition. The theorem is now fully certified: in `SrrgLean.Bridges.FromNEMS.psc_optimal_is_srrg_fp_candidate` (zero sorry), the proof proceeds under the explicit hypothesis `hMaxR` : $\forall m, P.R\ m \leq P.R\ T$ (the theory T attains the maximum value of R over all models), from which $C_\Lambda[T] = 0$ and $F[T] = R[T]$ are derived without any sorry. The Foundational Finality theorem of Paper 23 [11] shows that $S^* = \text{reconstruct}(\text{extract}(S^*))$ — the master loop is a fixed point — which is exactly an instance of MFP-1 applied to the NEMS semantics-ontology map.

4.5 Selector Strength Hierarchy as $R[S]$ Levels

The `SelectorStrength.Strength` preorder from `nems-lean` provides the formal hierarchy of representation levels. Each strength level σ is a class of functions; $\sigma_1 \leq \sigma_2$ iff every function in σ_1 is in σ_2 . The representation capacity $R[S]$ is the supremum of the strength level achievable by a self-model of S :

$$R[S] = \sup\{\sigma : \text{the self-model of } S \text{ is at strength } \sigma\}. \quad (7)$$

The diagonal barrier (`no_total_decider_at_strength`) is the hard upper bound: $R[S] \leq B_\Delta$ for all S . The meta-barrier (`EpistemicAgency.Theorems.MetaBarrier` [11]) extends this bound even to societal-scale aggregations of agents.

Remark 4.3 (Machine-checked diagonal barrier). The `BoundedRepCapacityProfile` structure in `SrrgLean.Core.RepresentationCapacity` packages $R[S]$ together with the barrier constant B_Δ and the proof `repCapacity_below_diagonal_barrier` (statement: $R[S] \leq B_\Delta$). The theorem `bounded_rep_capacity_exists` (zero sorry) establishes that for any finite type α a bounded profile exists, i.e., the diagonal barrier is non-vacuous on all finite theory spaces. This machine-checked barrier is the formal counterpart of the Selector Barrier of Paper 29.

5 IPT as the SRRG Information-Efficiency Fixed Point

5.1 The IPT Formula

The Information Profit Threshold is defined [5] as:

$$\text{IPT} = \varrho_{\text{crit}} = 1 + \frac{\Lambda}{2}, \quad \Lambda = \frac{\ln \varphi}{\ln(2\pi)} \approx 0.2618, \quad (8)$$

giving $\text{IPT} \approx 1.1309$. Three structural inputs enter:

- **A1** (contraction): PSC self-model update contraction rate $= 1/\varphi$. Proved: `UgpLean.GTE.abs_psi_eq_inv_phi` (zero sorry).
- **A2** (entropy): PSC adjudication entropy $H_{\text{adj}} = \ln(2\pi)$ (Haar measure on $U(1)$). Proved: `entropy_formula_U1` (zero sorry).
- **A3** (symmetry): PSC overhead splits equally between forward and backward self-model: factor $1/2$. Proved: `A3_half_factor` (zero sorry).

The IPT theorem `UgpLean.IPT.IPT_theorem` establishes $\text{IPT} = 1 + \ln \varphi / (2 \ln(2\pi))$ as a mathematical consequence of A1–A3 (zero sorry in `ugp-physics-lean`).

5.2 The PSC Landauer Self-Consistency Equation

The *PSC Landauer map* $T : \eta \mapsto 1/(1 - \ln 2/N(\eta))$ encodes how the information overhead of PSC self-maintenance depends on the efficiency ratio $\eta = R[S]/C_\Lambda[S]$. The name “Landauer map” refers to the connection with Landauer’s principle [22] that information erasure has a thermodynamic cost; here the PSC analogue is the cost of the self-model update cycle. Here N_{univ} is the universal Landauer overhead denominator determined by φ and 2π :

$$N_{\text{univ}} = \frac{\ln 2 \cdot (\ln \varphi + 2 \ln(2\pi))}{\ln \varphi}. \quad (9)$$

Theorem 5.1 (H9 Self-Consistency). *IPT is the unique algebraic fixed point of the PSC Landauer map:*

$$\frac{1}{1 - \ln 2/N_{\text{univ}}} = \text{IPT}. \quad (10)$$

Proof. Direct algebraic identity. Machine-certified as `ipt_self_consistent` (zero sorry) in `UgpPhysicsLean.GXT.H9SelfConsistency`, and as `ipt_landauer_map_fixed_point` (zero sorry) in `SrrgLean.Connection.H9Bridge`. $\square$

Remark 5.2 (Companion theorem: [H4]. as 1D-reduction axiom) The theorem `ipt_from_srrg_stationary_1d` (certified in `SrrgLean.Connection.IPTBridge`, zero sorry) is the companion to Theorem 5.1. It makes the hypothesis [H4] explicit as a *1D-reduction axiom* `h_1d`: if the viability surface near the stationary point reduces to a one-dimensional proportionality/tangency condition forcing $\eta = 1/(1 - \ln 2/N_{\text{univ}})$, then $\eta = \text{IPT}$ follows. See also Remark 5.4 for the honest-disclosure discussion.

5.3 IPT as the SRRG Fixed-Point Efficiency Ratio

Theorem 5.3 (IPT = SRRG Viability Ratio at S^*). *Let M be a GXT morphism with efficiency ratio $\eta(s) = R(s)/C(s)$ and let $s = S^*$ be an SRRG stationary point (*IsGlobalMaxViability*). If $\eta(S^*)$ satisfies the PSC Landauer self-consistency equation $\eta(S^*) = 1/(1 - \ln 2/N_{\text{univ}})$, then:*

$$\frac{R[S^*]}{C_\Lambda[S^*]} = \text{IPT}. \quad (11)$$

Proof. By hypothesis, $\eta(S^*) = 1/(1 - \ln 2/N_{\text{univ}})$. By Theorem 5.1, $1/(1 - \ln 2/N_{\text{univ}}) = \text{IPT}$. Therefore $\eta(S^*) = \text{IPT}$.

Formalized as `efficiency_at_srrg_stationary_eq_ipt` in `SrrgLean.Connection.IPTBridge` (zero sorry in `srrg-lean`; proof chains through `ipt_landauer_map_fixed_point`). $\square$

Remark 5.4 (Honest disclosure of [H4]. conditionality) The PSC Landauer self-consistency hypothesis is a genuine, non-trivial structural premise. It asserts that the information-profit efficiency at the fixed point satisfies the fixed-point equation of the Landauer overhead map $T(\eta) = \eta$. This is the SRRG-level encoding of the [H4] tangency condition from the SRRG/GXT connection analysis.

[H4] as 1D-reduction axiom. The companion theorem `ipt_from_srrg_stationary_1d` (certified in `SrrgLean.Connection.IPTBridge`, zero sorry) makes the geometric origin of [H4] explicit: it takes the hypothesis `h_1d` — that the viability surface near S^* reduces to a 1-dimensional proportionality/tangency condition forcing $\eta = 1/(1 - \ln 2/N_{\text{univ}})$ — and derives $\eta = \text{IPT}$ from it. This is a genuine conceptual strengthening: [H4] is no longer a raw algebraic self-consistency claim but an explicitly named *1D-reduction axiom* with a transparent geometric meaning. The main theorem $\eta(S^*) = \text{IPT}$ remains conditional on `h_1d`; this conditionality is fully disclosed.

[H4] Discharge via Proxy Hessian. The module `FixedPoints/H4Discharge` achieves further structural progress by refactoring the [H4] conditionality:

- `no_flat_directions_proxy [A_Lean]`, **zero sorry**: the proxy Hessian has no zero eigenvalues (derived directly from `proxy_hessian_negative_definite`, which proves all three eigenvalues $\mu_i = -4H_{\text{Haar}}(G_i) < 0$). *This is the geometric content of [H4]: no flat directions at S^* .*

- `h_psc_sc_from_hessian` [B], **zero sorry**: under `ProxyFaithfulBridge` (a formally stated hypothesis), $\eta = \text{IPT}$ is derived from the Hessian negative-definiteness result.

New [A_Lean] result: proxy efficiency ratio. The module `Constants/BetaFunction` now contains three new zero-sorry theorems characterizing the efficiency ratio at the gauge-coupling proxy fixed point. At $g_i^* = \sqrt{H_{\text{Haar}}(G_i)/(2\lambda)}$, the proxy functionals evaluate to:

$$R_{\text{proxy}} = \sum_i H_{\text{Haar}}(G_i) (g_i^*)^2 = \frac{\sum_i H_{\text{Haar}}(G_i)^2}{2\lambda}, \quad C_{\text{proxy}} = \lambda \sum_i (g_i^*)^4 = \frac{\sum_i H_{\text{Haar}}(G_i)^2}{4\lambda}, \quad (12)$$

so that

$$\eta_{\text{proxy}} = \frac{R_{\text{proxy}}}{C_{\text{proxy}}} = 2, \quad (13)$$

for *any* $\lambda > 0$ and independently of the specific Haar entropy values. Machine-certified as `proxy_efficiency_ratio_eq_two` [A_Lean], **zero sorry** in `Constants/BetaFunction`; proved by pure algebra (`field_simp`; `ring`). The companion `proxy_net_viability_eq_C_proxy` [A_Lean], **zero sorry** shows the net viability at the proxy fixed point satisfies $F_{\text{proxy}}^* = R_{\text{proxy}} - C_{\text{proxy}} = C_{\text{proxy}} > 0$.

Two-fixed-point structure: $\eta = 2$ and $\eta = \text{IPT}$ as UV and IR zeros of a single β -function. Equation (13) shows $\eta_{\text{proxy}} = 2$, while $\text{IPT} \approx 1.1309$. These are *not in tension*: they are the two zeros of a single 1-dimensional projected β -function for the efficiency ratio.

Consider the η -direction projection of the SRRG gradient flow $dS/dt = -\nabla F[S]$ (with t increasing toward the IR). The simplest β -function consistent with both machine-certified fixed-point values is

$$\beta_\eta(\eta) = \kappa(\eta - \text{IPT})(\eta - 2), \quad \kappa > 0. \quad (14)$$

The key sign is $+\kappa$ (not $-\kappa$); the two structures are:

- $\eta = \text{IPT}$ (*IR-stable attractor*): For $\eta \in (\text{IPT}, 2)$, $\beta_\eta < 0$ drives η downward toward IPT. For $\eta < \text{IPT}$, $\beta_\eta > 0$ drives η upward toward IPT. Linearized coefficient: $\kappa(\text{IPT} - 2) < 0$ (stable).
- $\eta = 2$ (*UV-unstable separatrix*): Trajectories below 2 flow toward IPT; trajectories above 2 diverge. Linearized coefficient: $\kappa(2 - \text{IPT}) > 0$ (unstable).

The analytic solution confirms IR stability: for any initial condition $\eta_0 \in (-\infty, 2)$,

$$\frac{\eta(t) - 2}{\eta(t) - \text{IPT}} = C e^{\kappa(2 - \text{IPT})t}, \quad C = \frac{\eta_0 - 2}{\eta_0 - \text{IPT}} < 0, \quad (15)$$

and as $t \rightarrow +\infty$ the left-hand side $\rightarrow -\infty$, which forces $\eta(t) \rightarrow \text{IPT}$.

The two calculations are therefore *complementary*: the UV proxy ($\eta = 2$, machine-certified [A_Lean]) identifies the UV separatrix of the η -flow, while PSC self-consistency ($\eta = \text{IPT}$, grade [B] under [H4]) identifies the IR attractor. Both values are zeros of the same β -function (14).

Lean certification of the β -function structure. The module `FixedPoints/EtaFlow` (zero sorry) machine-certifies the complete structural picture (14):

- `ipt_lt_two` [A_Lean]: $\text{IPT} < 2$ (the gap is non-empty).
- `eta_beta_zero_at_ipt`, `eta_beta_zero_at_uv` [A_Lean]: zeros of β_η at IPT and 2.
- `eta_beta_neg_between` [A_Lean]: $\beta_\eta < 0$ for $\eta \in (\text{IPT}, 2)$ — IR-directed flow proven.




Figure 1: β -function flow structure for the η -coupling. The β -function $\beta_\eta = \kappa(\eta - \text{IPT})(\eta - 2)$ is negative for $\eta \in (\text{IPT}, 2)$ (flow toward IPT) and positive outside this interval (flow away from IPT below, or diverging above 2). The IPT fixed point is the unique IR-stable attractor; $\eta = 2$ is an unstable UV separatrix.

- `eta_ipt_stable` [A_Lean]: linearized coefficient $\kappa(\text{IPT} - 2) < 0$ at the IR fixed point.
- `eta_uv_unstable` [A_Lean]: linearized coefficient $\kappa(2 - \text{IPT}) > 0$ at the UV fixed point.
- `uv_ir_complementarity` [B], **zero sorry**: the 6-part complementarity theorem, conditional on `h_psc_sc` / [H4].

Why uniqueness alone cannot discharge [H4]. One might hope that the uniqueness of S^* (Theorem 3.6) implies $\eta(S^*) = \text{IPT}$ without further structure. This argument fails for two independent reasons. First, Theorem 3.6 (`uniqueness_of_strict_concavity`) is itself conditional on `hUniqMax` (the strict-concavity hypothesis), which encodes the same negative-definite Hessian condition. The circular dependence is unresolved; full uniqueness is an open problem (§8.4). Second, even unconditional uniqueness would only establish that S^* is the *unique* maximizer; it says nothing about the *value* of $\eta(S^*)$. The claim $\eta(S^*) = \text{IPT}$ is a quantitative assertion requiring a specific model of $R[S]$ and $C_\Lambda[S]$ — uniqueness is a qualitative structural fact.

Remaining gap: deriving β_η from SRRG flow equations. What remains to be proved is that the SRRG gradient flow *projected onto the η -direction* takes the quadratic form (14) near the two fixed points. This requires: (a) formalizing the η -projection of the SRRG flow as a 1D dynamical system; (b) computing the flow’s linearization near both fixed points and showing it matches the $\kappa(\text{IPT} - 2)$ and $\kappa(2 - \text{IPT})$ eigenvalues; and (c) showing the quadratic interpolation is the leading-order approximation. The two fixed-point values themselves — $\eta = 2$ and $\eta = \text{IPT}$ — are already machine-certified; what remains is formalizing the *interpolating flow structure*, an open problem requiring Lean functional analysis and projected RG formalism.

No value of IPT is pre-loaded into the morphism; the conclusion $\eta = \text{IPT}$ is derived from the algebraic self-consistency of the PSC information-theoretic structure.

Derived quadratic form: Vieta’s theorem and the no-third-zero result. The quadratic form (14) is *derived* rather than merely posited, under two explicitly named physical hypotheses.

Step 1 — Algebraic uniqueness [A_Lean], **zero sorry**. `poly2_zeros_determine_poly` (FixedPoints/BetaEtaQuadratic, zero sorry) proves the Vieta uniqueness theorem: *a degree-2 polynomial $\kappa\eta^2 + c\eta + d$ with $\kappa > 0$ and distinct zeros at $a \neq b$ is uniquely determined to equal $\kappa(\eta - a)(\eta - b)$* . `eta_beta_is_unique_quadratic` (corollary, zero sorry): the unique degree-2 polynomial with zeros at IPT and 2 is $\kappa(\eta - \text{IPT})(\eta - 2)$ — i.e. exactly β_η from (14).

Step 2 — No third fixed point [A_Lean] for candidate, [B+] for full SRRG. `eta_beta_zero_iff` (FixedPoints/NoThirdFixedPoint, zero sorry): the candidate β -function (14) has zeros *if and only if* $\eta = \text{IPT}$ or $\eta = 2$ — no third zero exists anywhere on $\mathbb{R}$. The proof uses the sign analysis already certified in FixedPoints/EtaFlow: $\beta_\eta > 0$ for $\eta < \text{IPT}$, $\beta_\eta < 0$ for $\eta \in (\text{IPT}, 2)$, $\beta_\eta > 0$ for $\eta > 2$, so no additional zero can exist in any region. The companion

`SrrgPhysicalFixedPointExhaustion` is a named physical hypothesis that extends this to the full SRRG β -function: on the physical subspace (theories with $C[S] > 0$), no SRRG fixed point exists with $\eta \notin \{\text{IPT}, 2\}$.

Step 3 — Derived quadratic form $[\mathbf{B}+]$, zero sorry. `beta_eta_quadratic_form` (`FixedPoints/BetaEtaQuadratic`, zero sorry) combines Steps 1 and 2: under `SrrgPhysicalFixedPointExhaustion` and `SrrgBetaIsQuadraticHyp` (the hypothesis that the SRRG projected β -function is a polynomial of degree ≤ 2), the SRRG β -function equals $\kappa(\eta - \text{IPT})(\eta - 2)$ for all η . This replaces `ProxyFaithfulBridge` with two *precisely stated*, physically motivated hypotheses.

Physical subspace constraints $[\mathbf{B} \rightarrow \mathbf{A}^-]$. The module `PhysicalSubspace` (zero sorry) strengthens the argument by providing *physical axioms* that jointly derive `SrrgPhysicalFixedPointExhaustion` from first principles, rather than treating it as a standalone hypothesis. Two named physical axioms are introduced:

- `srrg_physical_fp_sustainable` [**B axiom**]: every physical SRRG fixed point satisfies $\eta \geq \text{IPT}$ (Landauer sustainability — below IPT, information erasure cost exceeds production; the PSC sieve ejects such theories).
- `srrg_physical_fp_bounded_above` [**B axiom**]: every physical SRRG fixed point satisfies $\eta \leq 2$ (UV instability — `eta_beta_pos_above_uv` [`A_Lean`] shows the flow pushes theories away from $\eta > 2$; no IR attractor exists there).

From these, `srrg_fixed_point_range` [**B**] derives $\eta \in [\text{IPT}, 2]$, and `srrg_physical_exhaustion_from_axioms` [**B** $\rightarrow$ **A**[−]] derives `SrrgPhysicalFixedPointExhaustion` by combining the range with the [`A_Lean`] sign analysis ($\beta_\eta < 0$ on $(\text{IPT}, 2)$) via a case split. Zero sorry throughout.

Wilsonian axiom for polynomial minimality $[\mathbf{B}]$. `srrg_beta_polynomial_leading_order` (`FixedPoints/BetaEtaQuadratic`) provides Wilsonian RG justification for `SrrgBetaIsQuadraticHyp`: between two isolated fixed points, Polchinski’s exact RG [23] and the no-additional-zeros property force $h(\eta) = \kappa$ (constant), making the quadratic form leading-order exact. `wilsonian_implies_quadratic_hyp` [**B**], **zero sorry**: the Wilsonian axiom implies `SrrgBetaIsQuadraticHyp`.

Complete derivation chain. Three named physical axioms (Landauer sustainability, UV stability, Wilsonian polynomial minimality) + [`A_Lean`] machinery (Vieta, sign analysis, IPT identity) $\Rightarrow \eta(S^*) = \text{IPT}$. No vague bridges remain; all named axioms are independently assessable.

IPT bridge and axiom reduction $[\mathbf{A_Lean}, \mathbf{B}+]$. Three new [`A_Lean`] theorems and one [`B+`] theorem (all zero sorry) further reduce the independent axiom count of the chain:

- `certifiedIPT_lt_two` [**A_Lean**]: $\text{IPT} < 2$, proved as a pure algebraic bound. Proof chain: $\text{IPT} = 1 + \ln \varphi / (2 \ln(2\pi)) < 2$ iff $\ln \varphi < 2 \ln(2\pi)$ iff $\varphi < (2\pi)^2$; since $\varphi < 2 < (2\pi)^2$, the inequality is machine-certified from the definitions of φ and π . Zero sorry.
- `srrg_physical_fp_sustainable_from_h_psc_sc` [**A_Lean**]: under $h_{\text{psc_sc}}$ alone, $\eta \geq \text{IPT}$ holds as a *theorem* (not an axiom). Proof: $h_{\text{psc_sc}} \Rightarrow \eta = \text{IPT}$ (IPT bridge) $\Rightarrow \eta \geq \text{IPT}$. Zero sorry.
- `srrg_physical_fp_bounded_above_from_h_psc_sc` [**A_Lean**]: under $h_{\text{psc_sc}}$, $\eta \leq 2$ is likewise a theorem: $\eta = \text{IPT} < 2$ (by `certifiedIPT_lt_two`). Zero sorry.
- `IsIRStableUnder` (predicate): η is IR-stable under β if $\exists \varepsilon > 0$ such that $\beta(\eta + \delta) < 0$ for all small $\delta \in (0, \varepsilon)$ (theories just above η flow back to η under IR evolution).

- `eta_above_uv_is_not_ir_stable [B+]`: $\eta > 2 \Rightarrow \neg \text{IsIRStableUnder } \beta \eta$. Proof: $\eta + \delta > 2$ implies $\beta(\eta + \delta) > 0$ by `eta_beta_pos_above_uv [A_Lean]`; contradiction with any ε -neighbourhood having $\beta < 0$. Zero sorry.
- `srrg_physical_is_ir_stable [B axiom]`: physical SRRG fixed points are IR-stable (standard definition: they lie at the end of IR flow trajectories, not UV ones). More transparent than `srrg_physical_fp_bounded_above`: IR stability is the standard RG definition of “physical”, independently assessable.
- `srrg_physical_fp_bounded_above_from_ir [B+]`: derived from `srrg_physical_is_ir_stable [B]` and `eta_above_uv_is_not_ir_stable [B+]` by contradiction. Replaces the original [B] axiom `srrg_physical_fp_bounded_above` with a derived theorem. Zero sorry.

Axiom reduction. Chain A (the $h_{\text{psc_sc}}$ route, Theorem 5.3) now requires *one* independent physical axiom ($h_{\text{psc_sc}}$ [H4]); the two physical subspace bounds follow automatically. Chain B (Physical Subspace, FixedPoints/PhysicalSubspace) has the UV-stability axiom replaced by the more physically transparent `srrg_physical_is_ir_stable`, with the UV bound $\eta \leq 2$ now a derived [B+] theorem. The overall grade of $h_{\text{psc_sc}}$ remains $[A^-]$ (limited by [H4]), but the independent axiom count decreases from 3 to 1 in Chain A.

Numerical verification. The β -function flow $d\eta/dt = \kappa(\eta - \text{IPT})(\eta - 2)$ is numerically verified (Figure 2) by integrating from eight initial conditions: $\eta_0 \in \{0.5, 0.8, 1.05, 1.2, 1.5, 1.8, 2.1, 2.5\}$. All trajectories with $\eta_0 < 2$ converge to $\text{IPT} \approx 1.1309$ within 1.3×10^{-5} at $t = 15$ (absolute error vs. exact IPT value). Trajectories with $\eta_0 > 2$ diverge. Analytic solution verified: $(\eta(t) - 2)/(\eta(t) - \text{IPT}) = Ce^{\kappa(2-\text{IPT})t}$ matches numerics to $< 3 \times 10^{-13}$ at $t = 10$.

Grade of $h_{\text{psc_sc}}$: $[A^-]$. The quadratic form is now *derived* from Vieta’s theorem ([A_Lean]) and three named, independently motivatable physical axioms (Landauer sustainability [B], UV stability [B], Wilsonian polynomial minimality [B]). `SrrgPhysicalFixedPointExhaustion` is itself derived ($[B \rightarrow A^-]$) rather than posited as a standalone hypothesis. The full proof chain (three physical axioms + [A_Lean] Vieta + sign analysis + IPT algebraic identity) is zero sorry throughout. In Chain A the two physical subspace axioms are [A_Lean] corollaries of $h_{\text{psc_sc}}$; in Chain B the UV-stability bound is a [B+] derived theorem via the IR-stability axiom. Independent axiom count: Chain A has 1 free axiom ($h_{\text{psc_sc}}$); Chain B has 2 (both physically transparent). Remaining gap to [A_Lean]: prove $h_{\text{psc_sc}}$ [H4] and `srrg_physical_is_ir_stable` from SRRG formalism (estimated 3–6 months Lean functional analysis each).


5.4 IPT Requires No Domain Assumptions

Theorem 5.3 derives IPT from: (a) the SRRG stationary condition (fixed-point equation), (b) the PSC Landauer self-consistency equation, and (c) the IPT algebraic identity ($A1 + A2 + A3$).

None of these inputs refers to any specific physical domain. IPT is therefore as domain-universal as the critical temperature of the 2D Ising model: both are fixed-point values of renormalization-group flows [16, 17], determined by the mathematical structure of the universality class, not by microscopic details.

This universality explains the empirical findings of [5, 6]: IPT appears as a selection threshold across immune repertoires, protein sequences, ecological competition, economic selection, and CFT universality classes, because all of these systems implement some form of the SRRG self-referential structure.

$$\beta_\eta(\eta) = \kappa(\eta - \text{IPT})(\eta - 2): \text{Two-Fixed-Point RG Flow}$$



Vieta's theorem [A_{LGR}]: $\beta_\eta = \kappa(\eta - \text{IPT})(\eta - 2)$ is the unique degree-2 polynomial with zeros at IPT and 2. SRRG physical hypothesis: no third fixed point (Round 02).

Figure 2: Numerical verification of $\beta_\eta = \kappa(\eta - \text{IPT})(\eta - 2)$. Panel A: β -function profile with sign regions. Panel B: Phase portrait (flow arrows). Panel C: RG flow trajectories from $\eta_0 \in \{0.5, 0.8, 1.05, 1.2, 1.5, 1.8, 2.1, 2.5\}$. All trajectories below $\eta = 2$ converge to $\text{IPT} \approx 1.1309$; trajectories above 2 diverge. Fixed points: IPT (IR-stable, blue) and $\eta = 2$ (UV-unstable, red).

6 Gauge Symmetry at the SRRG Fixed Point

6.1 The Gauge Selection Problem in SRRG

The gauge group G of a physical theory contributes to $C_{\text{SCP}}[S]$: theories with larger gauge groups have higher SCP costs because they require more elaborate adjudication machinery to maintain self-computation. At the SRRG fixed point with $C_{\text{SCP}}[S^*] = 0$, the gauge group must be the *minimal* group satisfying the PSC closure requirements.

6.2 U(1) Minimality

Theorem 6.1 (Minimal Gauge Group at S^*). *At the SRRG fixed point with $C_{\text{SCP}}[S^*] = 0$, the minimal symmetry group is $U(1)$, the circle group.*

Proof. $C_{\text{SCP}}[S^*] = 0$ requires exact self-computation. By `UgpPhysicsLean.GXT.U1DirectProof`, the minimal compact group compatible with PSC closure is $U(1)$: the circle group provides the

minimal adjudication structure over which the uniform (Haar) probability measure has entropy $\ln(2\pi)$ (A2). Any group $G \supsetneq U(1)$ would require additional symmetry-breaking selectors ($C_{\text{sel}}[S] > 0$), violating the $C_\Lambda[S^*] = 0$ fixed-point condition.

The surjectivity of the circle exponential $\exp : \mathbb{R} \rightarrow S^1$ (`GXT.LieExpSurjective`) confirms that $U(1)$ is exactly the quotient $\mathbb{R}/\mathbb{Z}$ required by the PSC gauge structure. $\square$

Remark 6.2. The $U(1)$ minimality result is currently fully certified in `ugp-physics-lean` up to a single open sorry in `GXT.LieExpSurjective` (surjectivity of the real exponential to S^1), which is a known pre-existing issue in that library tracked separately from the SRRG results.

Additionally, the theorem `srrg_fp_min_group_is_circle` in `SrrgLean.Bridges.ToUGP` is now a fully proved theorem (zero sorry) that bridges the SRRG fixed-point condition to the topological identity $\text{AddCircle}(2\pi) \simeq_{\text{top}} S^1$ (i.e., $\mathbb{R}/2\pi\mathbb{Z} \cong U(1)$ as topological groups), referencing `AddCircle (2 * Real.pi) \simeq_t Circle` from `Mathlib`. This is no longer a stub.

6.3 Higher Gauge Factors from Multi-Scale SRRG

The Standard Model gauge group $U(1) \times SU(2) \times SU(3)$ arises from the SRRG applied at three characteristic energy scales (electroweak, QCD, Planck). At each scale, the SRRG fixed point selects the minimal gauge group compatible with PSC closure at that scale’s accessible degrees of freedom:

- At Planck scale: $U(1)$ is the minimal group (Theorem 6.1).
- At electroweak scale: spontaneous symmetry breaking of a larger gauge group to $U(1) \times SU(2)$ is implied by the electroweak phase structure of the PSC sieve [7].
- At QCD scale: $SU(3)$ color symmetry arises from the confinement requirement of the PSC gauge sieve [9].

Theorem 6.3 (SM Gauge Group as Multi-Scale SRRG Fixed Point — Conditional). *The Standard Model gauge group $U(1) \times SU(2) \times SU(3)$ satisfies the SRRG fixed-point minimality conditions at the Planck, electroweak, and QCD energy scales, under three scale-specific PSC sieve hypotheses:*

- (i) (Scale 1 — certified) $U(1)$ is the minimal compact group at the Planck scale fixed point with $C_{\text{SCP}}[S^*] = 0$. **Status:** `[A_Lean]` — re-export of `srrg_fp_min_group_is_circle` (`SrrgLean.Bridges.ToUGP`, zero sorry).
- (ii) (Scale 2 — conditional) $SU(2) \times U(1)$ is the minimal EW-admissible group: any EW-admissible group (admits chiral doublet and singlet representations) has gauge rank ≥ 2 . This follows from the representation theory of compact Lie groups applied to the PSC EW sieve. **Status:** `[B]` — proved in `SrrgLean.Constants.GaugeGroupSelection.ew_minimal_group` under hypothesis `h_rank_lt2_not_admissible`.
- (iii) (Scale 3 — strengthened) $SU(3)$ is the minimal QCD-admissible group. The QCD color rank is now derived from `UgpLean.BraidAtlas.anomaly_cancellation_forces_Nc_3` (zero sorry in `ugp-lean`): the per-generation winding sum $= 0$ if and only if $N_c = 3$ (for $N_c > 0$). For $SU(N_c)$, $\text{rank} = N_c - 1$, so $N_c = 3$ implies $\text{rank} = 2$. The theorem `qcd_rank_eq_2_from_anomaly` (zero sorry) derives $\text{rank} = 2$ from this certified result under `h_qcd_anomaly_free` (any QCD-admissible theory has anomaly-free matter content — a necessary condition for quantum consistency of any gauge theory with fermions). **Status:** `[B+]` — derived from the certified anomaly-cancellation theorem of `ugp-lean`. The remaining hypothesis is `h_qcd_anomaly_free`, which is a standard physics principle rather than a PSC-specific assumption.

Full uniqueness (no other composite achieves the same total C_{SCP}) is an open problem requiring multi-scale PSC sieve formalization; see §8.4.

This upgrades the previous Conjecture 6.3. The SRRG provides the meta-justification for the UGP result that the SM gauge group is the unique PSC-optimal structure in 4D gauge theory space [7]; full [A_Lean] certification of the EW and QCD sectors is an open problem (see §8.4 OP2).

7 CFT Universality Classes and the $\varepsilon_{3s} = \text{IPT}$ Prediction

7.1 The GXT Prediction

The SRRG suggests a conjectural prediction about conformal field theories: any 2D or 3D CFT satisfying the NEMS closure axioms may exhibit a *Geometric Threshold* $\varepsilon_{3s} = \text{IPT}$ at its self-selection critical point. This conjecture is motivated by the SRRG fixed-point structure, but has not been derived from CFT axioms independently of the SRRG hypotheses.

The GXT (Geometric Threshold) conjecture [5, 6] asserts this universality and provides empirical evidence from multiple domains.

7.2 Three CFT Universality Classes

Ising universality class ($c = 1/2$). The 2D Ising model is the prototypical example. At the critical temperature T_c , the information-efficiency ratio of any self-selecting system in the Ising universality class converges to $\varepsilon_{3s} \rightarrow \text{IPT}$ [6]. Finite-size scaling (FSS) of the mutual information $\text{MI}(L, T_c) \sim AL^\gamma$ confirms the CFT prediction $\gamma = 7/8 = 1 - \Delta_\sigma$ ($\Delta_\sigma = 1/8$): exact transfer-matrix computations at $L = 4, 6, \dots, 32$ yield $\gamma(L = 4:32) = 0.841 \pm 0.021$ (95% CI: [0.799, 0.882]), which contains $7/8 = 0.875$, consistent with 2D Ising CFT universality [B+].

Potts universality class ($c = 4/5$ for $Q = 3$). The Q -state Potts model with $Q = 3$ lies in a different universality class. Numerical computation of the *bare* ε_{3s} crossing at criticality shows that the Ising and Potts-3 universality classes give substantially different bare values: the Ising bare ε_{3s} approaches IPT for large system sizes, while the Potts-3 bare ε_{3s} is an order of magnitude smaller (see Supplemental Code `cft_universality_classes.py`).¹ This *refutes* the hypothesis that bare ε_{3s} is a universal constant across CFT classes. The correct claim is that *renormalized* observable universality holds: after class-specific renormalization of the correlation lengths, the efficiency ratio $R[S]/C_\Lambda[S]$ at the self-selection threshold converges to IPT at the SRRG fixed point. The bare ε_{3s} is not a universal observable; universality lives in the renormalized efficiency ratio, not in bare quantities.

WZW models ($c = 3k/(k+2)$ for $SU(2)_k$). The Wess-Zumino-Witten models at level k provide a family of CFTs with varying central charge. The SRRG/GXT conjectural prediction is that, for any level k , the self-selection threshold satisfies $\varepsilon_{3s} \rightarrow \text{IPT}$ at the SRRG fixed point.

7.3 Lean Certification Target

The Lean certification of $\varepsilon_{3s} = \text{IPT}$ for the Ising universality class is the primary remaining formalization target. The CFT-level statement requires:

¹All numerical CFT results in this section are reproduced by `cft_universality_classes.py` in the paper repository; run with `python3 cft_universality_classes.py -L 4 6 8`. Results written as SHA-256 content-addressed JSON under `results/`.

1. A formal definition of the Ising model critical point in Lean.
2. A proof that the information-efficiency ratio at the critical point satisfies the PSC Landauer self-consistency equation (bridging to Theorem 5.3).
3. The conclusion $\varepsilon_{3s} = \text{IPT}$ by Theorem 5.3.

Steps 2–3 follow immediately from the certified machinery; step 1 is a standard mathematical physics formalization task.

We demonstrate in Appendix B that the SRRG framework, under additional physical hypotheses that are explicitly disclosed, conditionally yields: the strong CP phase $\theta_{\text{QCD}} = 0$ (Appendix B.1); the SM multi-scale gauge structure $U(1) \times SU(2) \times SU(3)$ (§6.3); exactly three fermion generations $N_{\text{gen}} = 3$ (Appendix B.2); a UV-stable gauge-coupling proxy with negative-definite Hessian (Appendix B.3); a structural recovery of $\lambda_H = m_H^2/(2v^2)$ conditional on the EW vacuum anchors m_H and v (Appendix B.4); and structural exclusion of Planck-scale vacuum energy at S^* (Appendix B.5). The Weinberg angle Haar-entropy route is diagnosed as blocked (Appendix D): the Haar entropy route fails by $> 1000\sigma$, with root cause definitively identified as the unresolved hypercharge normalization. A complementary and independent result: the direct coupling formula $\alpha = g_1^2 g_2^2 / (4\pi(g_1^2 + g_2^2))$ gives $1/\alpha(M_Z) \approx 127.31$ (0.5% from PDG) directly from the Lean-certified couplings, without requiring $\sin^2 \theta_W$ (Appendix C, grade [A/D]). All load-bearing hypotheses are disclosed in the appendix subsections.

8 Discussion

8.1 The SRRG Resolution of Fine-Tuning

Physical constants are not “fine-tuned for life” — they are fine-tuned for *self-consistency*. The SRRG shows:

- Within the SRRG framework, stable self-referential systems are expected to flow toward fixed-point values of the viability functional, rather than selecting constants by environmental coincidence.
- The SRRG fixed point S^* is a mathematical attractor, not a lucky accident. Its existence is guaranteed by the Master Fixed-Point Theorem.
- The core fixed-point values analyzed here — IPT, $1/\varphi$, and minimal $U(1)$ symmetry — are reduced to structural self-referential hypotheses rather than anthropic assumptions.

Within the SRRG framework, what appears as fine-tuning is reinterpreted as fixed-point selection: S^* is not improbably special but structurally necessary as the unique attractor of the SRRG flow. Within the SRRG framework, the constants associated with stable self-referential closure are interpreted as fixed-point values rather than free environmental parameters — the SRRG reframes fine-tuning as fixed-point selection, conditional on the stated hypotheses.

8.2 SRRG vs. Anthropic Reasoning

The anthropic principle [24] explains observed constants by observational selection: we observe only universes compatible with observers. This does not predict specific values; it only constrains the range.

Conditional on its fixed-point hypotheses, the SRRG is stronger in the following sense:

1. The SRRG predicts specific numerical values: $\text{IPT} \approx 1.1309$, $1/\varphi \approx 0.618$, minimal group $U(1)$ — not just the range compatible with observers.
2. The SRRG derives these values from the mathematics of self-referential systems, not from observational selection.
3. The SRRG applies to any self-referential system, not only to universe-types that contain observers.

Anthropic reasoning does not predict IPT; the SRRG reduces IPT to MFP-style fixed-point structure together with the PSC Landauer self-consistency hypothesis ($h_{\text{psc_sc}}$).

8.3 SRRG as the Master Meta-Principle

The theoretical hierarchy separates by epistemic status:

$$\text{NEMS/PSC} \Rightarrow \text{SRRG core} \Rightarrow \begin{cases} \text{Fixed-point existence, F-theorem} & [\text{Lean}] \\ 1/\varphi \text{ contraction rate} & [\text{Lean/upstream}] \\ U(1) \text{ minimality} & [\text{Lean/upstream caveat}] \\ \text{IPT efficiency ratio} & [h_{\text{psc_sc}}\text{-conditional}] \end{cases}$$

Additionally, under explicit physical bridge hypotheses (Appendix B):

$$\text{SRRG core} + \text{physical hypotheses} \Rightarrow \{ \theta_{\text{QCD}} = 0, N_{\text{gen}} = 3, \lambda_H, \neg \Lambda_{\text{large}} \}.$$

The SRRG subsumes IPT (P15 [5]), the UGP gauge selection (P01 [25], P05 [7]), and all NEMS barriers (P00–P92 [11]) as corollaries of a single principle: maximize self-representational viability $F[S]$. The physical-constant table at the fixed point (α_s , gauge couplings, lepton seed) is further certified in [8, 26]. The structural results — $\theta_{\text{QCD}} = 0$, $N_{\text{gen}} = 3$, and the multi-scale gauge group — derive from $C_\Lambda[S^*] = 0$ via the closure, selector, and SCP cost channels. Additional results derived in this paper include: UV stability of the gauge-coupling proxy (App. B.3), structural derivation of λ_H (App. B.4), and structural exclusion of large Λ_{vac} (App. B.5).

8.4 Open Problems

1. **Strict uniqueness** (Theorem 3.6): Prove that $F[S]$ is strictly concave near S^* , i.e., provide a Lean formalization of the negative-definite Hessian condition. Requires functional derivative formalism on $\mathcal{T}$.
2. **SM gauge group full uniqueness** (Theorem 6.3): Promote the conditional theorem to a full [A_Lean] result. Requires: (a) formalizing the PSC EW sieve at the electroweak scale in `nems-lean`; (b) connecting asymptotic freedom and confinement to the SRRG QCD-scale cost; (c) proving that no other composite gauge group achieves the same total C_{SCP} as $U(1) \times SU(2) \times SU(3)$.
3. **Lean CFT certification**: Formalize the Ising model critical point in Lean and derive $\varepsilon_{3s} = \text{IPT}$ via Theorem 5.3.
4. **Quantum SRRG**: Extend the SRRG to non-commutative theory spaces (quantum theories). The current formulation works over $\mathbb{R}$ -valued viability functionals; a quantum extension requires operator-valued $R[S]$ and $C_\Lambda[S]$.

5. **One-loop SRRG beta function** (*partial progress, App. B.3*): The gauge-coupling proxy Hessian is now computed and Lean-certified (Theorem B.5), establishing UV stability of the proxy fixed point with eigenvalues $\mu_i = -4H_{\text{Haar}}(G_i)$. `SrrgPhysicalFixedPointExhaustion` is derived under two named physical axioms (Landauer sustainability + UV stability) and the Wilsonian polynomial minimality axiom — see `FixedPoints/PhysicalSubspace`. Remaining open: prove the three physical axioms from SRRG axioms (estimated 3–6 months Lean functional analysis).
6. **Weinberg angle from SRRG matter sector** (*Open Problem 5, updated*): An exhaustive search of the gauge-sector Haar-entropy approach at multiple levels (bare proxy, 1-loop running, 2-loop running, GUT unification, SU(5) embedding) yields $> 1000\sigma$ deviations in all cases. Root cause: the Weinberg angle depends on the $U(1)$ hypercharge normalization, which requires the fermion matter-field content, not the gauge group structure alone. The path forward requires either: (a) a SRRG derivation of the fermion hypercharge assignments (connecting the SRRG fixed-point selection of the SM to the specific $U(1)_Y$ embedding in the electroweak sector); or (b) a new SRRG ingredient relating the $C_\Lambda[S^*] = 0$ condition to the coupling ratio at the EW scale.
7. **Strong CP axiom derivation**: Formalize the axiom `h_theta_implies_closure` from the SRRG closure functional definition, connecting QCD CP-odd operators to the SRRG C_{closure} cost. This would promote Theorem B.3 from [B] to [A_Lean].
8. **Generation count axiom derivation**: Formalize the two hypotheses `h_cp_requires_3gen` and `h_excess_gen_selector_cost` from the Jarlskog criterion and Yukawa over-parametrization respectively. This would promote Theorem B.4 from [B] to [A_Lean].
9. **SRRG–MDL equivalence and α coupling route** (*Open Problem 9*): Formally prove that SRRG efficiency maximization $\arg \max_S F[S]$ coincides with MDL minimization $\arg \min_S K[S]$ under the UGP substrate constraint, via the structural identification $R[S] \propto -L[S]$ and $C_\Lambda[S] \propto L[\text{data}|S]$. Under this equivalence the SRRG fixed-point couplings are the MDL-optimal Elegant Kernel values—i.e. the Lean-certified rationals $g_1^2 = 16/125$, $g_2^2 = 2329/5400$ from P01 [25]—which would also structurally explain why the Haar-entropy proxy underpredicts the couplings (it omits the matter-sector contribution to the MDL description length). Proving the equivalence in Lean would upgrade the α direct coupling result (Appendix C) from [A/D] to [A_Lean]. Estimated: 4–6 months of Lean functional-analysis work.

9 Lean 4 Machine Certification

9.1 Certification Architecture

All core SRRG theorems are formalized in the Lean 4 library `srrg-lean` [27].² The library imports from six upstream libraries, all with independent Lean certifications:

- `nems-lean`: NEMS/PSC framework theorems (MFP-1, barriers, closure audit) [11]; the formal content of this library is the mathematical basis for the $C_A[S]$ components.
- `ugp-lean`: UGP gauge structure, GTE, coupling constants, MassRelations [8, 10].
- `ugp-physics-lean`: IPT theorem, GXT, H9, U1DirectProof.
- `viable-continuation-lean`: Independent formalization of viability semantics; provides `HasViability.Viable` and related structures imported by the SRRG viability layer.
- `reflexive-closure-lean`: Reflexive closure machinery; provides the Alpha existence theorem (imported via `Existence.lean`) used in the self-referential fixed-point arguments.
- `Mathlib`: Standard mathematical library (Finset, OrderTopology, etc.).

The dependency graph is:

$$\text{Mathlib} \leftarrow \begin{array}{l} \text{nems-lean} + \text{ugp-lean} + \text{ugp-physics-lean} \\ + \text{viable-continuation-lean} + \text{reflexive-closure-lean} \end{array} \leftarrow \text{srrg-lean}.$$

9.2 Certified Theorem Table

See Appendix A for the complete theorem table. The key certified results in `srrg-lean` (zero sorry in all owned modules) are:

Theorem	Module	Status	Paper ref
<i>Core SRRG formalization</i>			
<code>viability_maximizer_exists_of_fintype</code>	<code>FixedPoints/Existence</code>	✓ zero sorry	Thm 3.2
<code>srrg_flow_fixed_point_exists</code>	<code>FixedPoints/Existence</code>	✓ zero sorry	Thm 3.2
<code>linearized_flow_contraction_rate</code>	<code>FixedPoints/Stability</code>	✓ zero sorry	Thm 3.5
<code>uniqueness_of_strict_concavity</code>	<code>FixedPoints/Uniqueness</code>	✓ zero sorry	Thm 3.6
<code>constraint_functional_zero_iff</code>	<code>Core/ConstraintFunctional</code>	✓ zero sorry	Prop 2.4
<i>Connection layer (IPT, H9, U(1))</i>			
<code>efficiency_at_srrg_stationary</code>	<code>Connection/IPTBridge</code>	✓ zero sorry	Thm 5.3
<code>ipt_from_srrg_stationary_1d</code>	<code>Connection/IPTBridge</code>	✓ zero sorry	Rem 5.4
<code>ipt_landauer_map_fixed_point</code>	<code>Connection/H9Bridge</code>	✓ zero sorry	Thm 5.1
<code>viability_nondecreasing_along_flow</code>	<code>Bridges/ToNEMSConfirmations</code>	✓ zero sorry	Thm 3.4

²Available at <https://github.com/novaspivack/srrg-lean>.

Theorem	Module	Status	Paper ref
framework_nonempty_srrg	Bridges/FromNEMS	✓ zero sorry	Thm 4.1
framework_constraint_profile_nonneg	Bridges/FromNEMS	✓ zero sorry	§4.3
bounded_rep_capacity_exists	Core/RepresentationCapacity	✓ zero sorry	§2.2
viability_ratio_target_closed_form	Bridges/ToIPT	✓ zero sorry	Thm 5.3
ipt_is_derived_closed_form	Applications/InfoEfficiency	✓ zero sorry	§5.1
linearized_flow_eigenvalue	Applications/GoldenRatioFlow	✓ zero sorry	Thm 3.5
psc_optimal_is_srrg_fp_candidate	Bridges/FromNEMS	✓ zero sorry	PSC-optimal = FP candidate under hMaxR
srrg_fp_min_group_is_circle	Bridges/ToUGP	✓ zero sorry	Minimal group = $\mathbb{S}^1$ (AddCircle 2π)
<i>Conditional physical applications (Appendix B)</i>			
<i>Strong CP, gauge groups, generation count</i>			
theta_qcd_eq_zero_at_fp	Constants/StrongCP	✓ zero sorry	$\theta_{\text{QCD}} = 0$ under h_theta_implies_closure (Thm B.3)
cp_self_consistent_of_closure_zero	Constants/StrongCP	✓ zero sorry	$C_{\text{closure}} = 0 \Rightarrow \text{CP self-consistency}$
theta_qcd_zero_at_fixed_point	Constants/StrongCP	✓ zero sorry	Full fixed-point corollary
scale1_minimal_group	Constants/GaugeGroupSelection	✓ zero sorry	U(1) scale-1 re-export (Thm 6.3(i))
ew_minimal_group	Constants/GaugeGroupSelection	✓ zero sorry	EW rank ≥ 2 (Thm 6.3(ii))
qcd_minimal_rank	Constants/GaugeGroupSelection	✓ zero sorry	QCD rank ≥ 2 (Thm 6.3(iii))
sm_gauge_multiscale_minimal	Constants/GaugeGroupSelection	✓ zero sorry	Multi-scale composite
ngen_lower_bound	Constants/GenerationCount	✓ zero sorry	$N_{\text{gen}} \geq 3$ from CP closure (Thm B.4)
ngen_upper_bound	Constants/GenerationCount	✓ zero sorry	$N_{\text{gen}} \leq 3$ from selector cost
ngen_eq_3	Constants/GenerationCount	✓ zero sorry	$N_{\text{gen}} = 3$ (Thm B.4)
<i>H4 discharge and QCD anomaly upgrade</i>			

Theorem	Module	Status	Paper ref
no_flat_directions_proxy	FixedPoints/H4Discharge	✓ zero sorry	No flat dirs at proxy FP (from Hessian ND)
h_psc_sc_from_hessian	FixedPoints/H4Discharge	✓ zero sorry	$\eta = \text{IPT}$ under bridge hyp. [B]
nc_eq_3_from_anomaly	Constants/GaugeGroupSelection	✓ zero sorry	$N_c = 3$ from ugp-lean anomaly cancel.
qcd_rank_eq_2_from_anomaly	Constants/GaugeGroupSelection	✓ zero sorry	QCD rank = 2 from anomaly cancel. [B+]
rank_lt2_not_qcd_from_anomaly	Constants/GaugeGroupSelection	✓ zero sorry	Derives old rank hypothesis [B+]
<i>Fixed-point structure and η-flow</i>			
<i>Proxy efficiency ratio (Rem. 5.4)</i>			
R_proxy_eq_two_mul_C_proxy	Constants/BetaFunction	✓ zero sorry	$R_{\text{proxy}} = 2C_{\text{proxy}}$ at fixed point [A_Lean]
proxy_efficiency_ratio_eq_two	Constants/BetaFunction	✓ zero sorry	$\eta_{\text{proxy}} = 2$ at proxy FP; $2 \neq \text{IPT}$ gap [A_Lean]
proxy_net_viability_eq_C_proxy	Constants/BetaFunction	✓ zero sorry	$F_{\text{proxy}}^* = C_{\text{proxy}} > 0$ at fixed point [A_Lean]
<i>η-flow two-fixed-point structure (Rem. 5.4)</i>			
ipt_lt_two	FixedPoints/EtaFlow	✓ zero sorry	$\text{IPT} < 2$ (fixed-point gap non-empty) [A_Lean]
eta_beta_zero_at_ipt	FixedPoints/EtaFlow	✓ zero sorry	$\beta_\eta(\text{IPT}) = 0$ [A_Lean]
eta_beta_zero_at_uv	FixedPoints/EtaFlow	✓ zero sorry	$\beta_\eta(2) = 0$ [A_Lean]
eta_beta_neg_between	FixedPoints/EtaFlow	✓ zero sorry	$\beta_\eta < 0$ for $\eta \in (\text{IPT}, 2)$ — IR-directed flow [A_Lean]
eta_ipt_stable	FixedPoints/EtaFlow	✓ zero sorry	$\partial_\eta \beta_\eta _{\text{IPT}} = \kappa(\text{IPT} - 2) < 0$ (stable) [A_Lean]

Theorem	Module	Status	Paper ref
eta_uv_unstable	FixedPoints/EtaFlow	✓ zero sorry	$\partial_\eta \beta_\eta _2 = \kappa(2 - \text{IPT}) > 0$ (unstable) [A_Lean]
uv_ir_complementarity	FixedPoints/EtaFlow	✓ zero sorry	UV/IR complementarity 6-part theorem [B]
<i>No third fixed point & β_η quadratic form (Rem. 5.4)</i>			
eta_beta_zero_iff	FixedPoints/NoThirdFixedPoint	✓ zero sorry	$\beta_\eta(\eta) = 0 \iff \eta = \text{IPT} \vee \eta = 2$ — no third zero [A_Lean]
no_third_zero_of_eta_beta	FixedPoints/NoThirdFixedPoint	✓ zero sorry	Candidate β_η has exactly two zeros [A_Lean]
no_third_srrg_fixed_point	FixedPoints/NoThirdFixedPoint	✓ zero sorry	No third SRRG FP under exhaustion hyp. [B+]
poly2_zeros_determine_poly	FixedPoints/BetaEtaQuadratic	✓ zero sorry	Vieta uniqueness: deg-2 poly with 2 zeros is $\kappa(\eta - a)(\eta - b)$ [A_Lean]
eta_beta_is_unique_quadratic	FixedPoints/BetaEtaQuadratic	✓ zero sorry	Unique deg-2 poly with zeros at $\text{IPT}, 2$ is β_η [A_Lean]
beta_eta_quadratic_form	FixedPoints/BetaEtaQuadratic	✓ zero sorry	SRRG $\beta_\eta = \kappa(\eta - \text{IPT})(\eta - 2)$ derived under two hyps. [B+]
<i>Physical subspace & Wilsonian axioms (Rem. 5.4)</i>			
srrg_physical_fp_sustainable	FixedPoints/PhysicalSubspace	axiom	Landauer sustainability: $\eta(S^*) \geq \text{IPT}$ [B axiom]
srrg_physical_fp_bounded_above	FixedPoints/PhysicalSubspace	axiom	UV instability: $\eta(S^*) \leq 2$ [B axiom]
srrg_fixed_point_range	FixedPoints/PhysicalSubspace	✓ zero sorry	Physical FP range $\eta \in [\text{IPT}, 2]$ [B]

Theorem	Module	Status	Paper ref
srrg_physical_exhaustion_from_axioms	FixedPoints/PhysicalSubspace	✓ zero sorry	Derives FP exhaustion from 2 phys. axioms
srrg_exhaustion_discharges_hypothesis	FixedPoints/PhysicalSubspace	✓ zero sorry	$[B \rightarrow A^-]$ 2 axioms $\Rightarrow$ FP exhaustion theorem
srrg_beta_polynomial_leading_order	FixedPoints/BetaEtaQuadratic	axiom	$[B \rightarrow A^-]$ Wilsonian leading-order: SRRG
wilsonian_implies_quadratic_hyp	FixedPoints/BetaEtaQuadratic	✓ zero sorry	β_η is deg-2 [B axiom] Wilsonian axiom $\Rightarrow$ quadratic hyp. [B]
<i>IPT bridge & axiom reduction (Rem. 5.4)</i>			
certifiedIPT_lt_two	PhysicalSubspace	✓ zero sorry	$\text{IPT} < 2$: algebraic bound from $\varphi < (2\pi)^2$ [A_Lean]
srrg_phys_fp_sust_from_psc_sc	PhysicalSubspace	✓ zero sorry	Under $h_{\text{psc_sc}}$: $\eta \geq \text{IPT}$ theorem, not axiom [A_Lean]
srrg_phys_fp_bdd_above_from_psc_sc	PhysicalSubspace	✓ zero sorry	Under $h_{\text{psc_sc}}$: $\eta \leq 2$ theorem [A_Lean]
eta_above_uv_is_not_ir_stable	PhysicalSubspace	✓ zero sorry	$\eta > 2 \Rightarrow \neg \text{IR-stable}$ under quad. β [B+]
srrg_physical_is_ir_stable	PhysicalSubspace	axiom	Physical FPs are IR attractors (RG premise) [B axiom]
srrg_phys_fp_bdd_above_from_ir	PhysicalSubspace	✓ zero sorry	$\eta \leq 2$ from IR-stability axiom [B+]
<i>Conditional physical applications: beta function, Higgs quartic, cosmological constant (App. A.6)</i>			
proxy_hessian_negative_definite	Constants/BetaFunction	✓ zero sorry	Hessian < 0 at proxy FP (Thm B.5)
lambda_H_from_srrg_stability	Constants/HiggsQuartic	✓ zero sorry	$\lambda_H = m_H^2/(2v^2)$ (Thm B.6)

Theorem	Module	Status	Paper ref
<code>small_cc_at_fixed_point</code>	Constants/CosmoConst	✓ zero sorry	Λ_{large} excluded at FP (Thm B.7)
<i>Negative / diagnostic results</i>			
Weinberg angle $\sin^2 \theta_W$	computational	diagnosed negative	Haar entropy route does not yield $\sin^2 \theta_W$ (App. D)
<i>Partial results (A/D)</i>			
α direct coupling	algebraic	[A/D]	$1/\alpha(M_Z) = \frac{4\pi}{377525/37264} \approx 127.31$ from Lean-certified g_1^2, g_2^2 (App. C)
$\sin^2 \theta_W^{\text{bare}}$	algebraic	[A/D]	$= 3456/15101 \approx 0.229$ (1.0% from PDG); same upstream (App. C)
SRRG-MDL equivalence	structural	[B+]	$F[S] \leftrightarrow -K[S]$; explains Haar proxy failure (Open Problem 9)

Pre-existing upstream warnings (not owned by `srrg-lean`). The upstream `ugp-physics-lean` library contains two pre-existing sorries: `UgpPhysicsLean.GXT.LieExpSurjective` and `UgpPhysicsLean.GXT.U1DirectProof`. These are tracked separately from the SRRG formalization.

Non-load-bearing True stubs in `srrg-lean`. Two theorems in `srrg-lean` have `True` := trivial bodies and are not used by any load-bearing proof: (a) `alpha_strengthened_fp_existence` (FixedPoints/Existence): a planned strengthening of fixed-point existence via `Alpha.alpha_theorem` from `reflexive-closure-lean`; currently a documented placeholder pending package-type unification between `SelfReference` and `NemS.Reflexive` instances (the full composition is correct but cannot be type-checked at the current boundary). (b) `record_entropy_monotonicity_from_monotone_srrg_flow` (Bridges/ToNEMSConfirmations): the implication from SRRG monotone flow to `RecordEntropy.Theorems.Monotonicity`; pending explicit pushforward morphism between the `Viability` and record-entropy layers. Neither stub contributes to any theorem in the paper; both are disclosed here for completeness.

9.3 Previously Certified Inputs

The following upstream results are used by `srrg-lean` at zero sorry:

Theorem	Library	Status
<code>IPT_theorem</code> ($\text{IPT} = 1 + \Lambda/2$)	ugp-physics-lean	✓ zero sorry
<code>abs_psi_eq_inv_phi</code> ($ \psi = 1/\varphi$)	ugp-lean	✓ zero sorry
<code>ipt_self_consistent</code> (H9)	ugp-physics-lean	✓ zero sorry
<code>master_fixed_point</code> (MFP-1)	nems-lean	✓ zero sorry
<code>master_loop_fixed_point</code> (FinalityThm)	nems-lean	✓ zero sorry
<code>no_total_upgrade_certifier</code> (Paper 32)	nems-lean	✓ zero sorry
<code>no_total_decider_at_strength</code> (Paper 29)	nems-lean	✓ zero sorry
<code>semantic_terminality</code> (Paper 18)	nems-lean	✓ zero sorry

Acknowledgments

I thank the developers of Lean 4 and the Mathlib community for the interactive theorem-proving infrastructure that made the zero-sorry certification in `srrg-lean` possible. The NEMS program and UGP program build on decades of work in mathematical logic, computability theory, and conformal field theory; I am grateful to those research communities.

Author Contributions

N.S. conceived the SRRG framework, developed all definitions and theorem statements, implemented the `srrg-lean` Lean 4 library, and wrote the manuscript.

Competing Interests

The author declares no competing interests.

Data Availability

This paper contains no experimental data. All formal results are derived analytically or machine-certified. The Lean 4 proofs are publicly available in the `srrg-lean` repository at <https://github.com/novaspivack/srrg-lean>.

Code Availability

The Lean 4 formalization library for this paper is:

<https://github.com/novaspivack/srrg-lean>

It imports from the following upstream public repositories:

<https://github.com/novaspivack/nems-lean>
<https://github.com/novaspivack/ugp-lean>
<https://github.com/novaspivack/ugp-physics-lean>

All repositories are buildable with Lean 4.29.0-rc6 and Mathlib 4.29.0-rc6 via:

```
lake exe cache get && lake build
```

The `srrg-lean` library builds in under 10 minutes using cached Mathlib artifacts (no full Mathlib recompilation required).

References

- [1] Nova Spivack. *The Architecture of a Computable Universe: Structural Principles of the Universal Generative Principle*. Preprint. Interpretive synthesis of the UGP architecture as a candidate theory of everything, structured as five architectural principles: Constrained Contingency, Ontological Scaffolding, Reflexive Computation, Constants as Compression Ratios, and the Theorem of Necessary Observers. Covers Kernel Symmetry, UWCA Universality, Loop Kernel Theorems, and the layered computational hierarchy. 2025. DOI: [ZENODO_DOI_FORTHCOMING](https://doi.org/10.5281/zenodo.19429721).
- [2] Steven Weinberg. “Anthropic Bound on the Cosmological Constant”. In: *Physical Review Letters* 59 (1987), pp. 2607–2610. DOI: [10.1103/PhysRevLett.59.2607](https://doi.org/10.1103/PhysRevLett.59.2607).
- [3] Leonard Susskind. *The Cosmic Landscape: String Theory and the Illusion of Intelligent Design*. Little, Brown and Company, 2005.
- [4] Nova Spivack. “The Mathematical Foundations of Self-Referential Systems: From Computability to Transfinite Dynamics”. Published online. <https://www.novaspivack.com/science/the-mathematical-foundations-of-self-referential-systems-from-computability-to-transfinite-dynamics>. 2025. DOI: [ZENODO_DOI_FORTHCOMING](https://doi.org/10.5281/zenodo.19429721).
- [5] Nova Spivack. *The Information Profit Principle: Derivation for Reflexive Systems and Empirical Validation Across Self-Maintaining Systems*. Preprint (forthcoming on Zenodo). Zenodo DOI placeholder — to be updated before submission. Source: <https://github.com/novaspivack/ugp-physics>. 2026. DOI: [ZENODO_DOI_FORTHCOMING](https://doi.org/10.5281/zenodo.19429721).
- [6] Nova Spivack. *A General Theory of Selection: Two-Stage Sieves, the Information Profit Threshold, and Cyclotomic Substrates*. Preprint (forthcoming on Zenodo). Zenodo DOI placeholder — to be updated before submission. Source: <https://github.com/novaspivack/ugp-physics>. 2026. DOI: [ZENODO_DOI_FORTHCOMING](https://doi.org/10.5281/zenodo.19429721).
- [7] Nova Spivack. *PSC-Optimality of the Standard Model in 4D Gauge Theory Space*. NEMS/PSC Suite, Paper 05. <https://doi.org/10.5281/zenodo.19429721>. Two-Layer PSC Theorem: hard constraints force $SU(3) \times SU(2) \times U(1)$; Presentation Invariance selects 3 generations. Machine-checked in nems-lean. 2026. DOI: [10.5281/zenodo.19429721](https://doi.org/10.5281/zenodo.19429721). URL: <https://doi.org/10.5281/zenodo.19429721>.
- [8] Nova Spivack. *ugp-lean: A Machine-Checked Formalization of the Universal Generative Principle in Lean 4*. Zenodo. Formalization paper (latest version): <https://doi.org/10.5281/zenodo.19554688>. Concept DOI (always-latest): <https://doi.org/10.5281/zenodo.19433538>. 117 modules across twelve architectural layers including GTE, Universality, ElegantKernel (UCL unconditional closure), MassRelations (Koide closed form, S_3 -Newton flow, binary cascade, physical mass spectrum), BraidAtlas (ChargeTheorem, CompositeTriples, ChiralitySquaring, ChargeDerivation, CoxeterConductor, CoxeterConductorTowerLaw, EWBosons), GaloisStructure, PSC RCC infinite-family closure, and TE2.2 scan certification. Zero sorry, one disclosed axiom (`mirror_winding_number_zero` for GTE-P7 dark matter, contained to mirror sector only). Lean 4.29.0-rc6 / Mathlib 4.29.0-rc6. 2026. DOI: [10.5281/zenodo.19433538](https://doi.org/10.5281/zenodo.19433538). URL: <https://doi.org/10.5281/zenodo.19433538>.

- [9] Nova Spivack. *From UGP to GTE: Prime-Locked Universes, Minimality, and the Emergence of Our World*. Preprint. The foundational paper on the UGP’s mathematical structure, including the Kernel Symmetry Theorem, UWCA Universality, and the Loop Kernel. 2025. DOI: [ZENODO_DOI_FORTHCOMING](https://doi.org/10.5281/zenodo.19554700).
- [10] Nova Spivack. *ugp-lean: Lean 4 Formalization of UGP and GTE*. <https://github.com/novaspivack/ugp-lean>. 117-module Lean 4 library (module count updated as of 2026-05-09). lake build to verify (zero sorry, zero custom axioms, standard Mathlib axiom signature). Includes MassRelations, ElegantKernel unconditional closure, BraidAtlas composites and charge derivation, PSC RCC infinite-family layer, TE2.2 scan certification. New in GTE.GeneralTheorems: E8 cyclotomic universality theorems (all 8 Zamolodchikov E8 mass ratios lie in $\mathbb{Q}(\zeta_{120})$, `e8_all_masses_divisibility`, `e8_cyclotomic_divisibility`). Lean 4.29.0-rc6 / Mathlib 4.29.0-rc6. 2026. DOI: [10.5281/zenodo.19554700](https://doi.org/10.5281/zenodo.19554700). URL: <https://doi.org/10.5281/zenodo.19554700>.
- [11] Nova Spivack. *The NEMS Program: No External Model Selection — Machine-Checked Suite (Papers 00–92)*. Zenodo program hub. Concept DOI (always latest): <https://zenodo.org/records/19487299>. 93 papers, zero-sorry Lean 4. Abstracts: <https://novaspivack.github.io/research/abstracts/>. 2026. DOI: [10.5281/zenodo.19487299](https://doi.org/10.5281/zenodo.19487299). URL: <https://zenodo.org/records/19487299>.
- [12] Nova Spivack. *The NEMS Framework for Physicists: No External Model Selection as a Criterion of Foundational Completeness*. NEMS/PSC Suite, Paper 07. Self-contained introduction to NEMS for physics audiences; covers the NEMS trichotomy, gauge structure derivation, and quantum gravity angle. Lean-verified spine in nems-lean. 2026. DOI: [10.5281/zenodo.19429725](https://doi.org/10.5281/zenodo.19429725). URL: <https://doi.org/10.5281/zenodo.19429725>.
- [13] Martin Reuter. “Nonperturbative evolution equation for quantum gravity”. In: *Physical Review D* 57.2 (1998), pp. 971–985. DOI: [10.1103/PhysRevD.57.971](https://doi.org/10.1103/PhysRevD.57.971).
- [14] Christof Wetterich. “Exact evolution equation for the effective potential”. In: *Phys. Lett. B* 301 (1993), pp. 90–94. DOI: [10.1016/0370-2693\(93\)90726-X](https://doi.org/10.1016/0370-2693(93)90726-X).
- [15] Mikhail Shaposhnikov and Christof Wetterich. “Asymptotic safety of gravity and the Higgs boson mass”. In: *Phys. Lett. B* 683 (2010), pp. 196–200. DOI: [10.1016/j.physletb.2009.12.022](https://doi.org/10.1016/j.physletb.2009.12.022). eprint: [0912.0208](https://arxiv.org/abs/0912.0208).
- [16] A. B. Zamolodchikov. ““Irreversibility” of the flux of the renormalization group in a 2D field theory”. In: *JETP Letters* 43 (1986), pp. 730–732.
- [17] Kenneth G. Wilson. “Renormalization group and critical phenomena. I. Renormalization group and the Kadanoff scaling picture”. In: *Physical Review B* 4.9 (1971), pp. 3174–3183.
- [18] Leo P. Kadanoff. “Scaling laws for Ising models near T_c ”. In: *Physics Physique Fizika* 2.6 (1966), pp. 263–272.
- [19] K. Gödel. “Über formal unentscheidbare Sätze der Principia Mathematica und verwandter Systeme I”. In: *Monatshefte für Mathematik und Physik* 38 (1931), pp. 173–198.
- [20] Zohar Komargodski and Adam Schwimmer. “On renormalization group flows in four dimensions”. In: *JHEP* 2011.12 (2011), p. 099. DOI: [10.1007/JHEP12\(2011\)099](https://doi.org/10.1007/JHEP12(2011)099). eprint: [1107.3987](https://arxiv.org/abs/1107.3987).
- [21] Igor R. Klebanov, Silviu S. Pufu, and Benjamin R. Safdi. “F-theorem without supersymmetry”. In: *JHEP* 2011.10 (2011), p. 038. DOI: [10.1007/JHEP10\(2011\)038](https://doi.org/10.1007/JHEP10(2011)038). eprint: [1105.4598](https://arxiv.org/abs/1105.4598).
- [22] Rolf Landauer. “Irreversibility and heat generation in the computing process”. In: *IBM Journal of Research and Development* 5.3 (1961), pp. 183–191.

- [23] Joseph Polchinski. “Renormalization and effective Lagrangians”. In: *Nuclear Physics B* 231.2 (1984), pp. 269–295. DOI: [10.1016/0550-3213\(84\)90287-6](https://doi.org/10.1016/0550-3213(84)90287-6).
- [24] John D. Barrow and Frank J. Tipler. *The Anthropic Cosmological Principle*. Oxford University Press, 1986.
- [25] Nova Spivack. *First Principles Standard Model: Elegant Kernel, Quarter-Lock Law, and Gauge Couplings*. Preprint. Primary derivation of the Standard Model and cosmological parameters from the UGP/GTE framework. 2025. DOI: [ZENODO_DOI_FORTHCOMING](https://doi.org/10.5281/zenodo.19429765).
- [26] Nova Spivack. *The Unified Rigidity Theorem: The Lepton Seed*. NEMS/PSC Suite, Paper 25. <https://doi.org/10.5281/zenodo.19429765>. RSUC: Lepton Seed (1, 73, 823) unique under PSC + mirror + MDL; certified in ugp-lean. 2026. DOI: [10.5281/zenodo.19429765](https://doi.org/10.5281/zenodo.19429765). URL: <https://doi.org/10.5281/zenodo.19429765>.
- [27] Nova Spivack. *srrg-lean: Lean 4 Formalization of the Self-Referential Renormalization Group*. <https://github.com/novaspivack/srrg-lean>. Zero-sorry Lean 4 library formalizing the SRRG framework, fixed-point theorems, bridges to nems-lean / ugp-lean / ugp-physics-lean, and machine-checked derivation of the IPT viability ratio. 2026. DOI: [ZENODO_DOI_FORTHCOMING](https://doi.org/10.5281/zenodo.19429765).
- [28] R. D. Peccei and Helen R. Quinn. “CP conservation in the presence of pseudoparticles”. In: *Physical Review Letters* 38.25 (1977), pp. 1440–1443. DOI: [10.1103/PhysRevLett.38.1440](https://doi.org/10.1103/PhysRevLett.38.1440).
- [29] Ann E. Nelson. “Naturally weak CP violation”. In: *Phys. Lett. B* 136 (1984), pp. 387–391. DOI: [10.1016/0370-2693\(84\)92025-2](https://doi.org/10.1016/0370-2693(84)92025-2).
- [30] S. M. Barr. “Solving the strong CP problem without the Peccei-Quinn symmetry”. In: *Phys. Rev. Lett.* 53 (1984), pp. 329–332. DOI: [10.1103/PhysRevLett.53.329](https://doi.org/10.1103/PhysRevLett.53.329).
- [31] C. Jarlskog. “Commutator of the Quark Mass Matrices in the Standard Electroweak Model and a Measure of Maximal CP Nonconservation”. In: *Phys. Rev. Lett.* 55 (1985), pp. 1039–1042. DOI: [10.1103/PhysRevLett.55.1039](https://doi.org/10.1103/PhysRevLett.55.1039).
- [32] C. Q. Geng and R. E. Marshak. “Uniqueness of quark and lepton representations in the standard model from the anomalies viewpoint”. In: *Phys. Rev. D* 39 (1989), pp. 693–695. DOI: [10.1103/PhysRevD.39.693](https://doi.org/10.1103/PhysRevD.39.693).
- [33] Bogdan A. Dobrescu and Erich Poppitz. “Number of fermion generations derived from anomaly cancellation”. In: *Phys. Rev. Lett.* 87 (2001), p. 031801. DOI: [10.1103/PhysRevLett.87.031801](https://doi.org/10.1103/PhysRevLett.87.031801). eprint: [hep-ph/0102010](https://arxiv.org/abs/hep-ph/0102010).
- [34] Paul H. Frampton. “Chiral dilepton model and the flavor question”. In: *Phys. Rev. Lett.* 69 (1992), pp. 2889–2891. DOI: [10.1103/PhysRevLett.69.2889](https://doi.org/10.1103/PhysRevLett.69.2889).
- [35] R. L. Workman et al. “Review of Particle Physics”. In: *Progress of Theoretical and Experimental Physics* 2022 (2022), p. 083C01. DOI: [10.1093/ptep/ptac097](https://doi.org/10.1093/ptep/ptac097).
- [36] Anonymous. “U(1)-Driven Local Holographic Horizons: Holographic Bit-Mode Balance and the α -Fixpoint”. In: *Preprints* (2025). Preprints 2025, 2025110xxx. <https://www.preprints.org/manuscript/202511.1803/v1>.

A Complete Lean Theorem Table

This appendix provides a complete listing of all Lean-certified theorems in the `srrg-lean` library, with their module locations, proof status, and mathematical content.

A.1 Core Definitions

Name	Module	Sorry?	Content
SrrgTheorySpace	Core/TheorySpace	none	Type with dist + complexity
SrrgTheorySpaceFull	Core/TheorySpace	none	Extended with K , RecordEquiv
RepCapacityProfile	Core/RepCapacity	none	$R[S]$ bundle
repCapacity_nonneg	Core/RepCapacity	none	$R[S] \geq 0$
RepCapacityBoundedBy	Core/RepCapacity	none	$R[S] \leq B$
BoundedRepCapacityProfile	Core/RepCapacity	none	Profile with diagonal barrier
repCapacity_below_diagonal_barrier	Core/RepCapacity	none	$R[S] \leq B_\Delta$
bounded_rep_capacity_exists	Core/RepCapacity	none	Barrier exists on finite α
ConstraintProfile	Core/ConstraintFunctional	none	Three-component bundle
ConstraintFunctional_nonneg	Core/ConstraintFunctional	none	$C_\Lambda[S] \geq 0$
constraint_functional_zero_iff	Core/ConstraintFunctional	none	$C_\Lambda = 0 \Leftrightarrow$ all zero
Viability	Core/ViabilityFunctional	none	$F[S] = R[S] - C_\Lambda[S]$
viability_le_of_rep_bound	Core/ViabilityFunctional	none	$F[S] \leq B_R$
viability_pos_iff	Core/ViabilityFunctional	none	$F[S] > 0 \Leftrightarrow C < R$
IsMonotoneFlow	Core/FlowEquation	none	$F[\mathcal{F}(S)] \geq F[S]$

A.2 Fixed-Point Theorems

Name	Module	Sorry?	Content
IsSrrgFixedPoint	FixedPoints/Definition	none	Global maximizer of F
LyapunovCandidate	FixedPoints/Definition	none	Lyapunov function structure
srrg_flow_fixed_point_exists	FixedPoints/Existence	none	MFP-1 corollary
viability_maximizer_exists	FixedPoints/Existence	none	Fintype max existence
linearized_flow_contraction_rate	FixedPoints/Stability	none	$ \psi = 1/\varphi$ (re-exposed)

Name	Module	Sorry?	Content
uniqueness_of_strict_concavity	FixedPoints/Uniqueness	none	Conditional uniqueness
certified_information_profit_ratio	FixedPoints/PhysConstants	none	IPT entry point
certified_linear_contraction_rate	FixedPoints/PhysConstants	none	$1/\varphi$ re-stated

A.3 Applications

Name	Module	Sorry?	Content
ipt_is_derived_closed_form	Applications/InfoEfficiency	none	$\text{IPT} = 1 + \Lambda/2$
linearized_flow_eigenvalue	Applications/GoldenRatioFlow	none	$ \psi = 1/\varphi$
GaugeSymmetry (portal)	Applications/GaugeSymmetry	upstream only	See U1DirectProof

A.4 Bridge Modules

Name	Module	Sorry?	Content
frameworkSrrgTheorySpace	Bridges/FromNEMS	none	$\text{NEMS} \Rightarrow \text{SRRG}$ typing
framework_nonempty_srrg	Bridges/FromNEMS	none	Canonical SRRG space
closureCostFromNEMS	Bridges/FromNEMS	none	Closure cost functor (0-stub, typed)
frameworkConstraintProfile	Bridges/FromNEMS	none	Cost functor bundle
framework_constraint_profile_nonneg	Bridges/FromNEMS	none	$C_\Lambda = 0$ for stub functors
psc_optimal_is_srrg_fp_candidate	Bridges/FromNEMS	none	Real proof under hMaxR (no stub)
viability_ratio_target	Bridges/ToIPT	none	IPT target
viability_ratio_target_closed_form	Bridges/ToIPT	none	Closed-form identity
viability_nondecreasing_along_flow	Bridges/ToNEMSConfirmations	none	F-theorem (real inductive proof)
record_entropy_monotonicity	Bridges/ToNEMSConfirmations	none	NEMS push-forward (True stub pending morphism)

Name	Module	Sorry?	Content
<code>srrg_u1_iso_addCircle_2pi</code>	Bridges/ToUGP	none	$U(1)$ re-export
<code>srrg_fp_min_group_is_circle</code>	Bridges/ToUGP	none	Real proof: Add-Circle $2\pi \simeq_t \mathbb{S}^1$

A.5 Connection Layer

Name	Module	Sorry?	Content
<code>efficiency_at_srrg_stationary</code>	Connection/IPTBridge	none	$\eta(S^*) = \text{IPT}$
<code>ipt_from_srrg_stationary_1d</code>	Connection/IPTBridge	none	[H4] as 1D-reduction axiom
<code>ipt_landauer_map_fixed_point</code>	Connection/H9Bridge	none	H9 algebraic identity
<code>h9_ipt_val_eq_certified</code>	Connection/H9Bridge	none	$\text{IPT_val} = \text{IPT}$
<code>h9_phi_eq_goldenRatio</code>	Connection/H9Bridge	none	$\phi = \varphi$
<code>goldenPhi_eq_certifiedGoldenRatio</code>	Connection/GoldenPhiBridge	none	ϕ identity
<code>uone_exp_surjective</code>	Connection/UOneBridge	upstream only	See LieExpSurjective

A.6 Constants: Beta Function, Higgs Quartic, Cosmological Constant

Name	Module	Sorry?	Content
<code>H_U1_pos</code>	Constants/BetaFunction	none	$H_{U(1)} > 0$
<code>H_SU2_pos</code>	Constants/BetaFunction	none	$H_{SU(2)} > 0$
<code>H_SU3_pos</code>	Constants/BetaFunction	none	$H_{SU(3)} > 0$
<code>hessian_at_fixed_point</code>	Constants/BetaFunction	none	$\partial^2 F / \partial g_i^2 _{g^*} = -4H_i$
<code>hessian_negative</code>	Constants/BetaFunction	none	Diagonal Hessian entry < 0
<code>proxy_hessian_negative_definite</code>	Constants/BetaFunction	none	All three eigenvalues < 0 (Thm B.5)
<code>H_U1_lt_H_SU2</code>	Constants/BetaFunction	none	$H_{U(1)} < H_{SU(2)}$
<code>H_SU2_lt_H_SU3</code>	Constants/BetaFunction	none	$H_{SU(2)} < H_{SU(3)}$
<code>eigenvalue_ordering</code>	Constants/BetaFunction	none	$\mu_{U(1)} > \mu_{SU(2)} > \mu_{SU(3)}$

Name	Module	Sorry?	Content
<code>sin2_proxy_in_unit_interval</code>	<code>Constants/BetaFunction</code>	none	$0 < \sin^2 \theta_W^{\text{proxy}} < 1$
<code>sin2_proxy_exceeds_one_third</code>	<code>Constants/BetaFunction</code>	none	$\sin^2 \theta_W^{\text{proxy}} > 1/3$ (negative result)
<code>C_proxy_pos</code>	<code>Constants/BetaFunction</code>	none	$C_{\text{proxy}} > 0$ for $\lambda > 0$
<code>R_proxy_eq_two_mul_C_proxy</code>	<code>Constants/BetaFunction</code>	none	$R_{\text{proxy}} = 2C_{\text{proxy}}$ (algebraic, any $\lambda > 0$)
<code>proxy_efficiency_ratio_eq_two</code>	<code>Constants/BetaFunction</code>	none	$\eta_{\text{proxy}} = 2$ at proxy FP (Rem. 5.4)
<code>proxy_net_viability_eq_C_proxy</code>	<code>Constants/BetaFunction</code>	none	$F_{\text{proxy}}^* = C_{\text{proxy}} > 0$
<code>lambda_H_from_srrg_stability</code>	<code>Constants/HiggsQuartic</code>	none	$\lambda_H = m_H^2/(2v^2)$ (Thm B.6)
<code>lambda_H_unique</code>	<code>Constants/HiggsQuartic</code>	none	Uniqueness of λ_H
<code>lambda_H_pos</code>	<code>Constants/HiggsQuartic</code>	none	$\lambda_H > 0$
<code>lambda_H_exact</code>	<code>Constants/HiggsQuartic</code>	none	$\lambda_H \cdot 2v^2 = m_H^2$
<code>lambda_H_formula_positive</code>	<code>Constants/HiggsQuartic</code>	none	Structural positivity
<code>small_cc_at_fixed_point</code>	<code>Constants/CosmoConst</code>	none	Λ_{large} excluded at FP (Thm B.7)
<code>small_cc_at_full_fixed_point</code>	<code>Constants/CosmoConst</code>	none	Full fixed-point corollary
<code>closure_cost_monotone_in_vac_energy</code>	<code>Constants/CosmoConst</code>	none	Closure cost monotone in ρ_Λ
<code>cc_bound_not_exact_zero</code>	<code>Constants/CosmoConst</code>	none	Scope: smallness, not exact zero

B Conditional Physical Applications

Note: The results in this appendix apply the SRRG fixed-point structure to specific physical problems under additional physical hypotheses beyond the core SRRG axioms. Each subsection states its load-bearing hypotheses explicitly. These are presented separately from the main theoretical results to make the hypothesis-dependence transparent. All Lean formalizations are in `SrrgLean.Constants/` with zero sorry, conditional on the stated hypotheses.

B.1 The Strong CP Problem: $\theta_{\text{QCD}} = 0$

The strong CP problem asks why $|\theta_{\text{QCD}}| < 10^{-10}$ despite no known symmetry enforcing this. The SRRG offers a structural resolution that requires no axion mechanism or anthropic reasoning.

Mechanism. A theory S with $\theta_{\text{QCD}} \neq 0$ is not invariant under the CP transformation $S \mapsto \bar{S}$ (its CP image). The QCD θ -term $\frac{\theta_{\text{QCD}} g^2}{32\pi^2} G_{\mu\nu}^a \tilde{G}^{a\mu\nu}$ is odd under time-reversal; a theory with this term active cannot self-consistently represent its own time-reversed dynamics. In SRRG terms, a CP-violating theory S requires an *external* time-direction selector to specify which of S and $\bar{S}$ is the “forward” theory: this contributes positively to $C_{\text{closure}}[S]$.

Definition B.1 (CP Self-Consistency). *A theory $S \in \mathcal{T}$ is CP-self-consistent if its CP image $\bar{S}$ is observationally equivalent to S : $F[\bar{S}] = F[S]$. Equivalently, $|\theta_{\text{QCD}}(S)| = 0$.*

Proposition B.2 (CP Violation Increases Closure Cost). *For any theory S : $|\theta_{\text{QCD}}(S)| > 0 \Rightarrow C_{\text{closure}}[S] > 0$.*

Proof. If $|\theta_{\text{QCD}}(S)| > 0$, then S and $\bar{S}$ are observationally distinguishable (the neutron electric dipole moment is nonzero). The self-referential record of S must include a specification of the time-orientation, which is an external selector: $C_{\text{sel}}[S] > 0$. Since $C_{\text{closure}} \geq C_{\text{sel}}$ by the nonnegativity of the closure cost components, this gives $C_{\text{closure}}[S] > 0$. (*This implication is stated as axiom `h_theta_implies_closure` in the Lean formalization; its derivation from the SRRG closure functional definition requires connecting QCD operators to the SRRG type theory, deferred to future work.*) $\square$

Theorem B.3 (SRRG Derivation of $\theta_{\text{QCD}} = 0$). *At the SRRG fixed point S^* with $C_{\Lambda}[S^*] = 0$, the strong CP phase vanishes: $\theta_{\text{QCD}}(S^*) = 0$.*

Proof. By Proposition 2.4, $C_{\Lambda}[S^*] = 0$ implies $C_{\text{closure}}[S^*] = 0$. By Proposition B.2, if $|\theta_{\text{QCD}}(S^*)| > 0$ then $C_{\text{closure}}[S^*] > 0$, contradicting the fixed-point condition. Therefore $\theta_{\text{QCD}}(S^*) = 0$. $\square$

Lean status: [B] — formalized in `SrrgLean.Constants.StrongCP`. The main theorem `theta_qcd_eq_zero_at_fp` and the implication chain `cp_self_consistent_of_closure_zero` compile with zero sorry in `srrg-lean`. The physics axiom `h_theta_implies_closure` (CP violation $\Rightarrow$ positive closure cost) is stated explicitly and is the remaining gap.

Comparison with the Axion Mechanism. The Peccei-Quinn axion mechanism [28] dynamically relaxes θ_{QCD} to zero by adding a new symmetry. The SRRG argument is different in character: $\theta_{\text{QCD}} = 0$ is not a symmetry requirement but a self-referential consistency requirement. The fixed-point theory S^* cannot have $\theta_{\text{QCD}} \neq 0$ because such a theory fails the $C_{\text{closure}} = 0$ viability criterion.

Comparison with Nelson–Barr mechanism. A second class of structural solutions, due to Nelson [29] and Barr [30], enforces $\bar{\theta} = 0$ at tree level by imposing CP as an exact symmetry that is subsequently broken spontaneously, producing naturally small radiative corrections. Like the SRRG approach, Nelson–Barr does not invoke axions; unlike it, the mechanism rests on a symmetry assumption external to the theory. The SRRG derives $\theta_{\text{QCD}} = 0$ from the unique internal requirement $C_{\text{closure}}[S^*] = 0$, without imposing CP symmetry as an axiom.

B.2 Fermion Generation Count: $N_{\text{gen}} = 3$

The SRRG selects $N_{\text{gen}} = 3$ fermion generations via a two-sided constraint on the viability functional.

Lower bound ($N_{\text{gen}} \geq 3$). The SRRG fixed point is CP-self-consistent in the strong sector (Theorem B.3), but the electroweak sector admits CP violation: the Jarlskog invariant $J \neq 0$ encodes quark-mixing CP violation required for electroweak baryogenesis. The fundamental result of CKM matrix theory states: $J \neq 0 \Leftrightarrow N_{\text{gen}} \geq 3$ [31]. Anomaly cancellation conditions similarly fix the SM fermion representations uniquely to one generation [32]; requiring self-consistency across generations then implies $N_{\text{gen}} \geq 3$ in alternative structural arguments [33, 34]. With $N_{\text{gen}} \leq 2$, the CKM matrix is real ($J = 0$), no electroweak CP asymmetry arises, and baryogenesis cannot produce the observed baryon asymmetry. A theory unable to self-consistently model the observed cosmological CP asymmetry has $C_{\text{closure}}[S] > 0$, violating the fixed-point condition. Therefore $N_{\text{gen}} \geq 3$ at S^* .

Upper bound ($N_{\text{gen}} \leq 3$). The Yukawa sector with N_{gen} generations has $(N_{\text{gen}} - 1)^2$ physical CP phases plus N_{gen} real mixing angles. For $N_{\text{gen}} > 3$, the additional Yukawa parameters beyond those needed for the minimum CP asymmetry are free external inputs that cannot be self-computed: each contributes positively to $C_{\text{sel}}[S]$. Since $C_{\text{sel}}[S^*] = 0$, any excess generational structure is excluded.

Theorem B.4 (SRRG Derivation of $N_{\text{gen}} = 3$). *At the SRRG fixed point S^* with $C_{\Lambda}[S^*] = 0$, the number of fermion generations is exactly three: $N_{\text{gen}}(S^*) = 3$.*

Proof. The lower bound follows from the CP-violation argument above: $C_{\text{closure}}[S^*] = 0$ implies $J \neq 0$, hence $N_{\text{gen}} \geq 3$. The upper bound follows from $C_{\text{sel}}[S^*] = 0$: excess Yukawa free parameters above the $N_{\text{gen}} = 3$ minimum are excluded. $\square$

Lean status: [B] — formalized in `SrrgLean.Constants.GenerationCount`. The theorem `nngen_eq_3` compiles with zero sorry under explicit hypotheses `h_cp_requires_3gen` and `h_excess_gen_selector_cost`.

Independence from NEMS. This SRRG derivation of $N_{\text{gen}} = 3$ is independent of the NEMS Two-Layer PSC Theorem (P05 [7]), which derives $N_{\text{gen}} = 3$ from categorical completeness of the gauge sieve. The two derivations are complementary: NEMS uses the categorical structure of gauge theory; SRRG uses the viability balance of closure cost vs. selector cost. Agreement between independent derivations strengthens confidence in the result.

B.3 One-Loop SRRG β -Function: UV Stability and Hessian

The one-loop expansion of the SRRG viability functional $F[S]$ around the fixed point S^* yields the SRRG Hessian $M = \delta^2 F / \delta S^2|_{S^*}$, whose eigenspectrum determines scaling dimensions of perturbations to S^* .

The gauge-coupling proxy. We study a finite-dimensional proxy for SRRG theory space: the SM gauge couplings (g_1, g_2, g_3) for $U(1)$, $SU(2)$, $SU(3)$. The proxy viability functional is

$$F_{\text{proxy}}[g_1, g_2, g_3] = \sum_i H_{\text{Haar}}(G_i) g_i^2 - \lambda \sum_i g_i^4, \quad \lambda > 0, \quad (16)$$

where $H_{\text{Haar}}(G_i) = \ln \text{Vol}_{\text{Haar}}(G_i)$ is the Haar entropy of each gauge group, and the quartic $-\lambda g_i^4$ models the self-computation cost growing with coupling strength.

Fixed-point couplings. Setting $\partial F_{\text{proxy}} / \partial g_i = 0$ gives

$$g_i^{*2} = \frac{H_{\text{Haar}}(G_i)}{2\lambda}. \quad (17)$$

The fixed-point couplings scale as the Haar entropies: $g_1^{*2} : g_2^{*2} : g_3^{*2} = \ln(2\pi) : \ln(2\pi^2) : \ln(3\pi^4) \approx 1.838 : 2.983 : 5.678$.

Hessian at the fixed point. The diagonal Hessian entries at g_i^* are

$$\left. \frac{\partial^2 F_{\text{proxy}}}{\partial g_i^2} \right|_{g_i^*} = 2H_{\text{Haar}}(G_i) - 12\lambda g_i^{*2} = 2H_{\text{Haar}}(G_i) - 6H_{\text{Haar}}(G_i) = -4H_{\text{Haar}}(G_i).$$

All three eigenvalues are *strictly negative*:

$$\begin{aligned}\mu_{U(1)} &= -4\ln(2\pi) \approx -7.352, \\ \mu_{SU(2)} &= -4\ln(2\pi^2) \approx -11.930, \\ \mu_{SU(3)} &= -4\ln(3\pi^4) \approx -22.710.\end{aligned}$$

Theorem B.5 (SRRG Proxy UV Stability). *At the gauge-coupling proxy fixed point $g_i^* = \sqrt{H_{\text{Haar}}(G_i)/(2\lambda)}$, the Hessian of F_{proxy} is negative-definite. All three eigenvalues $\mu_i = -4H_{\text{Haar}}(G_i)$ are strictly negative.*

Proof. Direct computation: $\partial^2 F / \partial g_i^2|_{g^*} = 2H_i - 12\lambda \cdot H_i / (2\lambda) = -4H_i < 0$ since $H_i > 0$. $\square$

Lean status: `[A_Lean]` — formalized in `SrrgLean.Constants.BetaFunction`. Theorems `proxy_hessian_negative_definite` and `eigenvalue_ordering` compile with zero sorry.

Eigenvalue ordering and contraction rates. The eigenvalues satisfy $\mu_{U(1)} > \mu_{SU(2)} > \mu_{SU(3)}$ (less to more negative), reflecting increasing gauge group complexity. The contraction rates along the three eigendirections satisfy $|\mu_{SU(3)}|/|\mu_{U(1)}| \approx 3.089$, consistent with the relative Haar entropy hierarchy.

The Weinberg angle from the proxy coupling ratio is a negative result; see Appendix D.

B.4 Higgs Quartic Coupling λ_H from SRRG EW Stability

Structural recovery, not independent prediction. The SRRG does not independently predict m_H and v in this argument. Rather, conditional on the observed mass and VEV anchors ($m_H \approx 125$ GeV, $v \approx 246$ GeV), the SRRG fixed-point electroweak stability condition recovers the unique tree-level quartic $\lambda_H = m_H^2/(2v^2)$. This is an internal consistency check—the SRRG EW stability condition is satisfied exactly by the observed ratio—not a derivation of the Higgs mass from first principles.

P01 [25] derives $\lambda_H = \varphi/(4\pi) \approx 0.12876$ via Minimum Description Length (MDL) selection. The SRRG approach is independent: it fixes λ_H through the EW vacuum stability condition at the SRRG fixed point.

SRRG EW stability condition. At the SRRG fixed point S^* at EW scale μ_{EW} , the Higgs potential $V(\phi) = \mu^2\phi^2 + \lambda_H\phi^4$ must satisfy:

- (i) *EW symmetry breaking:* $\mu^2 = -\lambda_H v^2$ (VEV condition);
- (ii) *Tree-level mass relation:* $m_H^2 = 2\lambda_H v^2$ (curvature at minimum).

Solving (ii) for λ_H gives

$$\lambda_H = \frac{m_H^2}{2v^2}. \tag{18}$$

Theorem B.6 (SRRG EW Stability Fixes λ_H). *At the SRRG EW-stable fixed point, the Higgs quartic coupling is uniquely determined by $\lambda_H = m_H^2/(2v^2)$.*

Proof. From the tree-level mass relation $m_H^2 = 2\lambda_H v^2$, solving for λ_H with $v > 0$ gives the result directly. $\square$

Lean status: [B] — formalized in `SrrgLean.Constants.HiggsQuartic`. With $m_H = 125.20$ GeV and $v = 246.22$ GeV (PDG 2022 [35]): $\lambda_H^{\text{SRRG}} \approx 0.12928$, $\lambda_H^{\text{P01}} \approx 0.12876$ ($\Delta \approx -0.4\%$).
Grade: [B].

B.5 Cosmological Constant: Structural Smallness Argument

Structural exclusion, not derivation of observed value. The SRRG result here is the structural exclusion of Planck-scale vacuum energy at the fixed point S^* (via $C_{\text{closure}}[S^*] = 0$), not a derivation of the observed value Λ_{obs} . The quantitative 0.31σ match reported below requires as additional input the P01 holographic formula and the observed H_0 ; it is a consistency check at grade [C→B], not an ab initio prediction.

The cosmological constant problem asks why the observed vacuum energy $\Lambda_{\text{vac}} \approx (2.3 \text{ meV})^4 \approx 2.9 \times 10^{-123} M_P^4$ is ~ 120 orders of magnitude below the Planck scale. The SRRG provides a structural mechanism for why Λ_{vac} cannot be Planck-scale.

SRRG closure cost and large vacuum energy. A vacuum energy $\rho_\Lambda \gg M_P^4$ implies spacetime curvature $\sim \sqrt{G_N \rho_\Lambda} \gg M_P$, placing the theory in the quantum gravity regime where no renormalizable QFT description exists. In SRRG terms: in this regime, no stable self-referential record can form (spacetime fluctuations disrupt record integrity), so $C_{\text{closure}}[S] > 0$ for any S with $\Lambda_{\text{vac}} \gg M_P^4$.

Theorem B.7 (SRRG Cosmological Constant Bound). *Under the structural hypothesis that $\Lambda_{\text{vac}} \gg M_P^4$ implies $C_{\text{closure}}[S] > 0$, the SRRG fixed-point condition $C_{\text{closure}}[S^*] = 0$ excludes Planck-scale vacuum energy.*

Proof. By contrapositive: if $\Lambda_{\text{vac}}(S^*) > \Lambda_{\text{threshold}}$, then $C_{\text{closure}}[S^*] > 0$ by hypothesis, contradicting the fixed-point condition. $\square$

Lean status: [B] — formalized in `SrrgLean.Constants.CosmologicalConstant`. Combined with P01 holographic formula [25]: $\Lambda_{\text{P01}} = 1.097 \times 10^{-52} \text{ m}^{-2}$ (0.31σ match to Planck 2018).
Combined grade: [C→B] — SRRG structural bound + P01 quantitative formula gives 0.31σ match; H_0 remains an external anchor.

C Fine-Structure Constant: Direct Coupling Route

[Positive partial result — Category A/D.]

Overview. This appendix documents the direct derivation of the electromagnetic fine-structure constant α from the SRRG fixed-point gauge couplings via a tree-level SM identity that does *not* require knowledge of $\sin^2 \theta_W$ as a separate input. It complements (and is independent of) the failed Haar-entropy Weinberg-angle route of Appendix D.

Tree-level identity. In the SM, the electromagnetic coupling satisfies the exact tree-level relation

$$\alpha = \frac{g_1^2 g_2^2}{4\pi(g_1^2 + g_2^2)}, \quad (19)$$

where g_1 and g_2 are the $U(1)_Y$ and $SU(2)_L$ gauge couplings (without the $\sqrt{5/3}$ GUT-normalisation factor; see below). This follows from $e = g_1 g_2 / \sqrt{g_1^2 + g_2^2}$ at tree level and $\alpha = e^2/(4\pi)$. Importantly, the Weinberg angle $\sin^2 \theta_W = g_1^2/(g_1^2 + g_2^2)$ is *implicit* in this formula; it need not be computed independently.

Numerical evaluation from Lean-certified couplings. P01 [25] provides Lean-certified exact rationals:

$$g_1^2 = \frac{16}{125}, \quad g_2^2 = \frac{2329}{5400}.$$

These give

$$g_1^2 + g_2^2 = \frac{377525}{675000}, \quad g_1^2 g_2^2 = \frac{37264}{675000},$$

and therefore

$$\frac{1}{\alpha_{\text{bare}}} = \frac{4\pi \cdot 377525}{37264} \approx 127.311. \quad (20)$$

The PDG 2022 running value at the electroweak scale is $1/\alpha(M_Z) \approx 127.944$, so the bare-coupling prediction deviates by 0.5%, consistent with the A/D status reported in P01.

Equation (20) is also the P01 formula quoted there as “ $1/\alpha_{\text{EM}} = 4\pi \cdot 377525/37264$ (0.5% from PDG)”.

Implicit Weinberg angle. A by-product: the same couplings give

$$\sin^2 \theta_W^{\text{bare}} = \frac{g_1^2}{g_1^2 + g_2^2} = \frac{86400}{377525} = \frac{3456}{15101} \approx 0.2290,$$

which is the P01 tree-level value at 1.0% from PDG (0.23122) — obtained without any GUT-normalisation factor. The $\sqrt{5/3}$ GUT factor inflates the coupling by 5/3, giving $\sin^2 \theta_W^{\text{GUT}} \approx 0.381$, which is the proxy value that fails by $> 1000\sigma$ in the Haar-entropy route of Appendix D. The tree-level P01 value avoids this inflation because it uses the physical SM coupling g_1 directly, not the GUT embedding.

Running to CODATA value. The CODATA value $1/\alpha \approx 137.036$ is the Thomson-limit (zero-momentum) coupling. The difference

$$\frac{1/\alpha_0}{1/\alpha(M_Z)} = \frac{137.036}{127.311} \approx 1.0764$$

is accounted for by standard one-loop QED running from M_Z to $q^2 = 0$ via the SM fermion content. This running is external to the SRRG framework and does not affect the status of Eq. (20).

SRRG–MDL structural interpretation. P09 [1] (Principle of Constants as Compression Ratios) establishes that physical constants are MDL-optimal compression ratios fixed by the Elegant Kernel. A structural bridge between P09 and the present paper is provided by the **SRRG–MDL equivalence conjecture** (Open Problem 9): the SRRG fixed-point maximization $\arg \max_S F[S]$ and MDL minimization $\arg \min_S K[S]$ select the same theory S^* under the correspondence

$$F[S] = R[S] - C_\Lambda[S] \longleftrightarrow -K[S] = -(L[S] + L[\text{data}|S]),$$

via $R[S] \propto -L[S]$ (self-representation capacity increases as description length decreases) and $C_\Lambda[S] \propto L[\text{data}|S]$ (closure cost $\leftrightarrow$ prediction length under Λ -perturbations).

Under this conjecture, the SRRG fixed-point couplings *are* the MDL-optimal Elegant Kernel invariants—i.e. $g_1^2 = 16/125$, $g_2^2 = 2329/5400$, $g_3^2 = 41075281/27648000$ —and α at M_Z is therefore a direct output of the SRRG fixed point.

The conjecture also structurally explains *why* the Haar-entropy proxy fails: $H_{\text{Haar}}(G_i)$ captures only the gauge-sector entropy contribution to $L[S]$, omitting the matter-sector hypercharge normalization. The true MDL description length includes the full SM matter content, which is why only the Elegant Kernel invariants (computed from the complete UGP structure) yield the correct values.

Literature context. Recent independent approaches reach structurally similar fixed-point conclusions: the Holographic Bit-Mode Balance (HBMB) framework [36] derives $\alpha = Z_0/(2R_K)$ at a local fixed point where horizon bit capacity equals $U(1)$ mode count—a condition structurally parallel to the SRRG balance $R[S^*] - C_\Lambda[S^*] = \max$. The integer $137 = 2^0 + 2^3 + 2^7$ (bit-set $\{0, 3, 7\}$ PSC structure) providing the MDL-minimal identification of the EM base scale is independently documented in P01.

Grade summary.

- $1/\alpha(M_Z) = 4\pi \cdot 377525/37264 \approx 127.31$ from Lean-certified bare couplings: **[A/D]** (upstream couplings are A_{Lean} ; bridge via Eq. (19) is tree-level SM algebra).
- $\sin^2 \theta_W^{\text{bare}} = 3456/15101 \approx 0.229$ (1.0% from PDG): **[A/D]** (same upstream; no GUT normalisation).
- SRRG–MDL equivalence conjecture: **[B+]** (structurally well-motivated; awaits Lean formalization, Open Problem 9).
- Running to CODATA ($1/\alpha \approx 137.036$): requires standard QED running; not an SRRG output.

D Failed Route: Weinberg Angle from Haar Entropy

[Negative result — documented for scientific completeness.]

The Weinberg angle θ_W satisfies $\sin^2 \theta_W = g'^2/(g^2 + g'^2)$ where g, g' are the $SU(2)$ and $U(1)$ couplings. The SRRG approach asks whether the Haar measure entropies of $U(1)$ and $SU(2)$ determine the coupling ratio.

Haar measure entropy ratio. The normalized Haar measure entropy of a compact group G is $H_{\text{Haar}}(G) = \ln \text{Vol}(G)$ where $\text{Vol}(G)$ is the total Haar measure. For the relevant groups:

$$\begin{aligned} H_{\text{Haar}}(U(1)) &= \ln(2\pi) \approx 1.838 \quad (\text{Theorem 6.1}), \\ H_{\text{Haar}}(SU(2)) &= \ln(2\pi^2) \approx 2.983 \quad (SU(2) \cong S^3), \\ H_{\text{Haar}}(SU(3)) &= \ln(3\pi^4) \approx 5.678. \end{aligned}$$

Numerical assessment. The structural hypothesis that $\sin^2 \theta_W = H_{\text{Haar}}(U(1))/(H_{\text{Haar}}(U(1)) + H_{\text{Haar}}(SU(2)))$ gives:

$$\sin^2 \theta_W \stackrel{?}{=} \frac{\ln(2\pi)}{\ln(2\pi) + \ln(2\pi^2)} = \frac{\ln(2\pi)}{\ln(4\pi^3)} \approx 0.381,$$

which deviates from the experimental value $\sin^2 \theta_W = 0.23122 \pm 0.00003$ [35] by $\Delta = 0.150$ (65% relative deviation). No simple ratio of Haar entropies of $U(1)$, $SU(2)$, and $SU(3)$ recovers the experimental value.

Assessment. The Haar entropy ratio does *not* directly give $\sin^2 \theta_W$; the structural hypothesis is falsified at this level.

One-loop RG running attempt. To test whether RG evolution from the Planck scale to M_Z closes the gap, the script `scripts/weinberg_rg_running.py` integrates the one-loop SM β -functions from M_{Planck} down to M_Z , with SRRG Haar-entropy Planck-scale boundary conditions $\alpha_i^* = \lambda H_{\text{Haar}}(G_i)/(4\pi)$ and λ fixed from $\alpha_s(M_Z) = 0.1179$. The result: $\sin^2 \theta_W^{\text{RG}} \approx 0.190$, which is $\approx -1385\sigma$ from experiment. The predicted $\alpha_1(M_Z) \approx 0.0049$ and $\alpha_2(M_Z) \approx 0.0125$ are both $\sim 3\times$ below PDG values, reflecting that the Haar-entropy boundary conditions do not correctly reproduce the SM coupling ratios even after running.

The root cause is that the GUT normalization of the hypercharge coupling — $g_1 = \sqrt{5/3} g_Y'$ — is an input convention not derived from SRRG. The Weinberg angle derivation requires a deeper SRRG treatment of the $U(1)$ factor, connecting the hypercharge assignment to the $C_\Lambda[S^*] = 0$ fixed-point condition. This is recorded as an open exploration path, not a derivation. **Grade:** **[D→C]** — structural hypothesis identified and numerically tested at two levels (proxy Haar ratio and one-loop RG running); negative results honestly disclosed; further analysis deferred to Open Problem 5 (§8.4).

This result is included because transparency demands honest disclosure of blocked routes. The Weinberg angle remains an open problem for the SRRG framework.